

HiGrEC: A Controlled Study of Structural Inductive Bias for Response Clarity Classification in Political Interviews

Vidvathama R Siddarth Gottumukkula Sanjana Reddy Vonteri Shashikanta Sahoo
2024122002 2023102040 2026901007 2026900007

Team Nier_ANLP
International Institute of Information Technology, Hyderabad

Abstract

SemEval-2026 Task 6 (CLARITY) asks how a politician’s answer responds to an interview question: one of nine evasion types (Subtask 2), which fixes one of three clarity levels (Subtask 1). We report a controlled, pre-registered study of what limits single-pass classifiers on this task. A DeBERTa-v3-large classifier, improved one measured change at a time, reaches 0.384 multi-reference macro-F1 on Subtask 2 per model and 0.405 as a ten-seed system with a one-parameter logit adjustment. Every attempt to improve the decision rather than the model (re-ranking, three hierarchies, rebalancing losses, model soups) failed, and error analysis points to the model, not label noise, as the bottleneck. Replacing only the backbone with Qwen3-8B fine-tuned with LoRA, with data, input, loss and selection unchanged, improves Subtask 2 by 0.092 and Subtask 1 by 0.095 on all ten paired seeds and reaches 0.543 / 0.746 as a system (+0.136, 95% CI [0.054, 0.224]). The gain is spread across classes rather than concentrated on the label boundaries we predicted. A longer schedule is adopted on a held-out slice of the training data; training on all data adds no distinguishable gain. All scores are on the development set, since test labels were not released.

1 Introduction

Politicians often answer the question they would have liked to be asked. Thomas et al. (2024) turned this into a classification task with a two-level taxonomy of response clarity and evasion, and SemEval-2026 Task 6 (Thomas et al., 2026) made it a shared task: Subtask 2 asks for one of nine evasion types, and Subtask 1 for one of three clarity levels that the nine types determine.

The competition was won by multi-call large language model (LLM) pipelines (Shi et al., 2026; Peerachaidecho and Rutherford, 2026), while fine-tuned encoders stopped at about 0.50 macro-F1

on Subtask 2 regardless of scale or ensembling (Thomas et al., 2026). A leaderboard does not say *why* encoders stop there. The candidate explanations make different predictions: the labels may be too noisy to learn (then annotators would disagree where models fail), the decision rule may be wrong for the metric (then a better rule would help), the model may lack the knowledge needed to read an answer against its question (then a larger pretrained model would help under the same protocol), or the label meanings may not be learnable from 3,448 examples (then supplying definitions would help).

We test these explanations with a deliberately controlled single-pass classifier (Figure 1): every change is made one at a time against a named control, on ten paired seeds, with predictions registered before each run and no choice made on the development set. Our contributions are:

- A DeBERTa-v3-large system (§4) whose two real improvements (training length and the full journalist question) are separated from documented negative results: re-ranking, hierarchical models, rebalancing losses, per-class thresholds and model soups.
- Evidence that the remaining gap lies in the model rather than in label noise or the decision rule (§6).
- A one-change test of the knowledge explanation (§5): replacing the backbone with Qwen3-8B-Base fine-tuned with LoRA (Qwen Team, 2025; Hu et al., 2022) improves both subtasks on all ten seeds and raises the ten-seed system from 0.405 to 0.543 (Subtask 2) and from 0.648 to 0.746 (Subtask 1).
- An analysis of where that gain falls, which refutes our own registered prediction (§6).

Test labels were never released, so all comparisons are on the 308-item development set, with seed




Figure 1: The classification system. One backbone (DeBERTa-v3-large, fully fine-tuned, or Qwen3-8B-Base with LoRA) reads the sub-question, the journalist’s full question and the answer, and a new linear head gives a distribution over the nine evasion types. A system averages ten seeds and applies logit adjustment, which divides by the train class prior π_c raised to τ ; the Subtask 1 clarity level follows from the Subtask 2 label through the fixed taxonomy.

Evasion type	Clarity level	Train %	Dev sets
Explicit	Clear Reply	30.5	115
Implicit	Ambivalent	14.2	99
Dodging	Ambivalent	20.5	96
General	Ambivalent	11.2	113
Deflection	Ambivalent	11.0	51
Partial/half-answer	Ambivalent	2.3	14
Declining to answer	Clear Non-Reply	4.2	17
Claims ignorance	Clear Non-Reply	3.5	15
Clarification	Clear Non-Reply	2.7	4

Table 1: The taxonomy: each evasion type determines a clarity level. Train %: share of the 3,448 training labels. Dev sets: number of the 308 dev items whose reference set contains the type.

spreads, paired differences and bootstrap intervals throughout.

2 Problem Statement, Data and Evaluation

Problem. SemEval-2026 Task 6 (Thomas et al., 2026) uses the question–answer pairs of Thomas et al. (2024), taken from U.S. presidential interviews. An item is a triple $x = (q, Q, a)$: a sub-question q (median 15 tokens), the journalist’s full turn Q from which it was split (median 62 tokens), and the president’s entire answer a (median 266 tokens). Let $\mathcal{Y}_9$ be the nine evasion types and $\mathcal{Y}_3$ the three clarity levels, and let $g : \mathcal{Y}_9 \rightarrow \mathcal{Y}_3$ be the fixed taxonomy map of Table 1. Subtask 2 (S2) asks for $f(x) \in \mathcal{Y}_9$, the way that sub-question was answered; Subtask 1 (S1) asks for a level in $\mathcal{Y}_3$. Like the top-ranked systems (Shi et al., 2026; Peerachaidecho and Rutherford, 2026), we learn f and predict $g(f(x))$ for S1. A classifier is *single-pass* if it makes one forward pass of one model per item; the question we ask is what limits such a classifier on this task.

Data. Train has 3,448 items with one label each. Dev has 308 items labelled by three annotators: 125 are unanimous, 150 split 2–1 and 33 three ways, so the reference set G_i of an item holds one to three types. Test has 237 items with two annotators each; its labels were never released and the evaluation server has closed, so every comparison in this paper is on dev. Two properties shape the modelling. First, 69% of training rows share their answer with another sub-question, and 71.7% of those shared answers carry different labels, so the model must attribute parts of a long answer to the right sub-question. Second, answers are long: for 101 of the 308 dev items the sub-question and answer exceed 512 DeBERTa tokens.

Metric. Both subtasks are scored by macro-F1 over all declared classes; a class that is never predicted contributes $F_1 = 0$. S1 uses plain macro-F1. S2 uses *multi-reference macro-F1*: a prediction $\hat{y}_i$ is a true positive for class $\hat{y}_i$ if $\hat{y}_i \in G_i$, a false positive if not, and a miss charges one false negative to every member of G_i . If every prediction is acceptable, every class predicted at least once has $F_1 = 1$, and

$$\text{macro-F1} = \frac{|\{\hat{y}_1, \dots, \hat{y}_N\}|}{9}. \quad (1)$$

The score therefore rewards two things: landing in the reference set (the *in-set rate*) and *coverage*, naming each class at least once, so Clarification (4 dev reference sets) counts as much as Explicit (115). The floors show the trade-off: always predicting Explicit lands in-set for 37.3% of dev items but scores 0.060 on S2, uniform guessing scores 0.129, and TF-IDF with logistic regression 0.257 (S1: 0.136, 0.299 and 0.445).

Evaluation protocol. *Epoch selection* uses the *train slice*, a fixed stratified 10% of train (345

rows) held out from training; dev is never used to choose anything. Every system claim rests on *10 paired seeds* (the same seed numbers in both arms), reported as mean $\pm$ standard deviation with the paired difference, its t statistic and the number of seeds that improve; variants are first screened on three seeds against a bar fixed in advance (+0.015 with at least two of three seeds up). A *single model* is one seed with argmax and no tuning on dev. A *system* averages the seeds’ probabilities and applies logit adjustment (Menon et al., 2021):

$$\hat{y}(x) = \arg \max_c \frac{\bar{p}(c | x)}{\pi_c^\tau}, \quad (2)$$

where $\bar{p}$ is the mean of the seeds’ probabilities, π_c the frequency of class c in train, and τ the only fitted parameter. Because τ is fitted on dev, it is chosen by 5-fold nested cross-validation (grid -1 to 2 , step 0.05): each fold’s τ maximises S2 macro-F1 on the other four, and the pooled held-out predictions are scored; we report the first cross-validation split and, where given, the mean over ten. Two systems are compared by a bootstrap over dev items with the seeds fixed (2,000 resamples), and because test has two annotators we also rescore dev against its three two-annotator reference sets. The hypotheses and predicted outcomes of every experiment were recorded in an experiment log before it ran.

3 Related Work

Task and top systems. Thomas et al. (2024) introduced the clarity and evasion taxonomy over U.S. presidential interviews and found that predicting the nine-way label and mapping it up the hierarchy beat predicting clarity directly. In the shared task (Thomas et al., 2026), the organisers’ fine-tuned Llama-70B scored 0.82 (S1) and 0.57 (S2) on test; LLM prompting and hierarchical use of the taxonomy were most effective, and fine-tuned encoders levelled off near 0.50 on test S2. TeleAI (Shi et al., 2026) ranked first (test 0.89 S1, 0.68 S2) with three confidence-routed chain-of-thought stages on DeepSeek-V3 whose prompts include boundary examples mined from the training confusion matrix; it averages 1.94 calls per item and reaches 0.617 (S2) and 0.812 (S1) on dev, and in the same paper a directly fine-tuned Qwen2.5-7B reached 0.495 S2 on dev. ChulaNLP (Peerachaidecho and Rutherford, 2026), second on S2, passed the top-five labels of a fine-tuned DeBERTa-large to few-shot Kimi-K2, scoring 0.52 S2 on dev against 0.46 for the encoder

alone. As test labels are unreleased, we compare with these systems on dev.

Other participating systems. PFW (Tamsal and Rusert, 2026) averaged the logits of 50 DeBERTa models per architecture (5 folds $\times$ 10 seeds) and reached 0.76 (S1) and 0.50 (S2) on test; three post-hoc techniques that raised their out-of-fold scores lowered their test scores by 0.02–0.10, much as per-class decision weights overfit dev for us. SilkPeak (Tetakali and Tetakali, 2026) trained one DeBERTa-v3-large with joint clarity and evasion heads (0.76 S1 on test) and report it ahead of hierarchical bi-encoders and a LoRA-tuned LLaMA-3-8B. CLaC (Turk et al., 2026) found prompted LLMs ahead of fine-tuned encoders (third on S2, 0.59); ttda704 (Tran Tan and Dinh, 2026) found structured chain-of-thought prompting (0.515 S2) ahead of parameter-efficient fine-tuning of Qwen3 models; and Team Evaluators (Nuthakki et al., 2026) reached 0.54 S2 with LoRA-tuned open LLMs. Each of these systems changes several components at once; we change one at a time, asking what limits a fine-tuned encoder and how much of the gap to multi-call pipelines one fine-tuned LLM closes.

Methods we build on. Logit adjustment (Menon et al., 2021) subtracts scaled log class priors from the logits after training; Balanced Softmax (Ren et al., 2020) corrects the prior inside the loss, and focal loss (Lin et al., 2017) down-weights easy examples. Our multi-reference analysis treats annotator disagreement as signal (Uma et al., 2021). Our encoder is DeBERTa-v3-large (He et al., 2023); our LLM classifier is Qwen3-8B-Base (Qwen Team, 2025) fine-tuned with LoRA (Hu et al., 2022) as a single-pass classifier.

4 The Encoder Track

Model and training. We fine-tuned all 435M parameters of DeBERTa-v3-large (He et al., 2023) as a 9-way classifier: the [CLS] state passes through a dense tanh pooler and a new linear layer over the nine labels. The baseline input pairs the sub-question (at most 96 tokens) with the answer, cut from the right at 512 tokens. The final input puts “Sub-question: ... Full question: ...” first, with up to 256 tokens, and raises the context to 1,024 tokens: 69% of training rows share their answer with another sub-question, and a model that cannot see the other sub-questions cannot tell which part

Configuration	Single model		System	
	S2	S1	S2	S1
Baseline, 8 ep.	0.315 ± 0.040	0.576 ± 0.030	0.428	0.601
+ 16 epochs	0.362 ± 0.026	0.601 ± 0.016	0.400	0.620
+ full question	0.384 ± 0.030	0.614 ± 0.027	0.405	0.648
All of train [†]	0.416 ± 0.020	0.614 ± 0.016	0.489	–

Table 2: DeBERTa-v3-large on dev. Single model: mean $\pm$ sd over seeds 0–9. System: 10-seed mean probability with logit adjustment (τ by nested CV). Rows 2–3 each add one change; row 3 is the final system. [†]All of train, last epoch, seeds 0–4 only (a 5-seed system, not comparable).




Figure 2: Single DeBERTa-v3-large models on dev for the three configurations: (a) S2 multi-reference macro-F1, (b) S1 macro-F1. Open circles joined by thin lines are individual seeds, paired across configurations; filled squares are the mean over the 10 seeds, with error bars of ± 1 standard deviation.

of the answer responds to which. Training used plain cross-entropy, AdamW at 10^{-5} with layer-wise decay 0.95, a cosine schedule and batch size 16 (Appendix A); the epoch was chosen on the train slice. A seed fixes the head initialisation and data order, so seeds 0–9 are paired across configurations.

What improved the model. Two changes improved single models (Table 2, Figure 2). Training for 16 epochs instead of 8 raised S2 by +0.047 ($t = 3.4$, 9/10 seeds up) and S1 by +0.026 ($t = 2.7$, 8/10). Adding the full question gave a further +0.022 S2 ($t = 2.0$, 7/10) and +0.012 S1 ($t = 1.3$, 6/10); together, +0.069 S2 and +0.038 S1, each on 9 of 10 seeds. Most of the gain comes from longer training, which also lifted the classes the baseline handled worst: General rose from 0.119 to 0.295 per model (seeds 0–4).

Undertraining at first looked like a negative result. The full question at 8 epochs scored below

the baseline (0.282 ± 0.043 against 0.337 ± 0.044 , seeds 0–4), but its final training loss (1.32–1.50, against 0.57–1.09) showed that it had not finished fitting: its seeds were less confident (mean top probability 0.445 against 0.594) and leaned on the class prior. At 16 epochs the same input gave the best single model, so every run now reports its final training loss and selected epoch next to its scores.

From model to system. Better single models did not give a better S2 system. Logit adjustment lifted the baseline ensemble by +0.109 (to 0.428) and the final ensemble not at all (0.409 to 0.405); the S2 difference of -0.022 between the two systems has a 95% bootstrap interval of $[-0.101, +0.061]$, while on S1 the final system is ahead (0.648 against 0.601). Our reading (inferred, not tested) is that longer training and logit adjustment correct the same fault, the pull towards frequent classes. Five-seed systems are unreliable on 308 items: the baseline system scored 0.438 on seeds 0–4 and 0.356 on seeds 5–9, a spread larger than most effects we wanted to measure, so all system claims use ten seeds. Training on all of train for 16 epochs and keeping the last epoch gained +0.029 S2 per model on seeds 0–4 (4/5 up) and lost 0.017 S1; as a 5-seed result it is not a ranking.

What did not help. Seven further ideas, each tested against a named control, left the system where it was. A second encoder that re-ranked the baseline’s top-3 labels from their definitions scored 0.354 against 0.365 for the ensemble it re-ranked: learning what each label means from 3,448 examples is harder than classifying them. Balanced Softmax with focal loss (Ren et al., 2020; Lin et al., 2017) over-corrected towards rare classes (Explicit F1 fell from 0.679 to 0.319) and used up the gain that logit adjustment would have given. Three hierarchy designs (routing through the clarity level, a factorised head, and a Non-Reply gate with branch specialists) did not beat the flat softmax, which already scores all nine labels jointly, and boundary experts for the three most-confused pairs changed S2 by +0.002, because they seldom overturn a decision the flat model makes from the same input. A uniform model soup of the ten models scored 0.267, because seeds that differ in head initialisation and data order do not average; nine per-class decision weights overfit 308 items, and an ensemble mixing both inputs scored 0.403 against 0.405. Encoders or rules trained on the same 3,448 examples did not move the system; §5 tests the backbone itself.

5 Is the Gap Knowledge? An 8B LLM Classifier

The encoder track and the error analysis (§6) show that the model, not label noise, is what fails. Two explanations remain: the encoder may lack *knowledge* that a 0.4B-parameter model has not acquired, or *label meaning*, which a classifier must infer from about 3,400 examples. A larger pretrained classifier tests the first with nothing else changed. Before any run we registered three predictions: the gain passes the screening bar; the single-model S2 mean lies in 0.43–0.50 and S1 rises by +0.02 to +0.05; and the gain falls mostly on the commitment boundary (Explicit $\leftrightarrow$ Implicit, General $\leftrightarrow$ Implicit/Explicit) and more on agreed than on contested items.

One change: the backbone. We took the best encoder configuration (sub-question, full question and answer, 16 epochs) and replaced DeBERTa-v3-large with Qwen3-8B-Base (Qwen Team, 2025), adapted with LoRA (Hu et al., 2022) of rank 16 ($\alpha = 32$, dropout 0.05) on all linear projections of the attention and feed-forward blocks. A new 9-way linear head, trained in full, reads the hidden state of the last token (Figure 1). The input text and budgets are those of the encoder, with one fixed closing question appended (“How does the answer respond to the sub-question?”), so the token the head reads sits after the whole question and answer in every item. Sequences are right-padded, and each run first checks that an item’s logits are the same alone and inside a padded batch, aborting otherwise; this guards against a head that reads a padding token. Training ran in bf16 with AdamW at 10^{-4} , a cosine schedule with 5% warm-up, 3 epochs and an effective batch of 16; the learning rate and epoch count change because the backbone does. Everything else is identical: the same 3,103 training rows, the same 345-row train slice (asserted identical in code), plain cross-entropy and the epoch with the best slice macro-F1. Seeds 0–9 are paired with the DeBERTa runs of the same seeds; Appendix A lists all hyperparameters and the compute.

Results at ten paired seeds. The 3-seed screen passed the bar (+0.081 on S2, 3 of 3 seeds up), so the run was extended to ten seeds (Table 3). Per model, Qwen improves S2 by +0.092 ($t = 7.12$) and S1 by +0.095 ($t = 8.48$) on 10 of 10 seeds, more than the whole encoder track gained (0.315 to 0.384). As a system it reaches 0.543 / 0.746

against 0.405 / 0.648, and the ranking holds on the 2-annotator reference sets (0.524 against 0.382). All registered size predictions held (the S1 gain exceeded its range); the prediction about *where* the gain falls was refuted (§6). Among published dev results (Figure 3), the system lies above TeleAI’s fine-tuned Qwen2.5-7B (0.495) and ChulaNLP’s pipeline (0.52), and below TeleAI’s multi-call pipeline (0.617) and the human reference (0.684); these come from other protocols, so the comparison locates our system rather than ranking it.

Two single-change follow-ups. In the 3-epoch runs the slice F1 rose at every epoch and the last epoch was always selected; under a cosine schedule this proves nothing, because the schedule itself favours the last epoch, so only a longer schedule tests undertraining. With 12 epochs (seeds 0–2) the slice selected epochs 9, 11 and 8 and the slice F1 rose by +0.084 on 3 of 3 seeds, so the rule registered in advance adopted the longer schedule; training loss fell to about 0.02 without the slice F1 declining, and dev S2 was 0.543 ± 0.015 (reported only). Training on all 3,448 rows gave +0.017 per model (6 of 10 seeds up), but the system difference has a 95% interval of $[-0.026, +0.092]$ and S1 fell from 0.746 to 0.711; as with the encoder, the 3-seed screen (+0.039) overstated the gain.

6 Analysis

The error analysis uses the final DeBERTa system’s 10-seed ensemble with argmax; the comparison uses the paired 10-seed runs of §5. All figures are on dev.

The model, not label noise, fails. Of the ensemble’s 308 predictions, 137 (44.5%) fell outside the item’s reference set, and 53 of these were on items where all three annotators agreed. Annotator disagreement is real and informative (Uma et al., 2021), but it does not account for the gap: scored as the test set is scored, each annotator against the other two reaches 0.643–0.766 on S2 (mean 0.684), while the ensemble reaches 0.380–0.381.

Commitment, not only topic. The largest error pairs (reference $\rightarrow$ prediction) are Implicit $\rightarrow$ Explicit (25), General $\rightarrow$ Implicit (20) and General $\rightarrow$ Explicit (20): they concern how committed an answer is rather than whether it stays on topic. Length also matters: the 20 inputs still cut at 1,024 tokens are in-set only 20% of the time (0.580 for the

Dev macro-F1	Qwen3-8B + LoRA	DeBERTa-v3-large	Δ	Seeds up	95% CI
Single model, S2	0.476 $\pm$ 0.043	0.384 $\pm$ 0.030	+0.092	10/10	
Single model, S1	0.709 $\pm$ 0.031	0.614 $\pm$ 0.027	+0.095	10/10	
Single model + LA, S2	0.538 $\pm$ 0.043	0.389 $\pm$ 0.019	+0.149	10/10	
System, S2	0.543 (0.549)	0.405 (0.412)	+0.136		[+0.054, +0.224]
System, S1	0.746	0.648	+0.098		
System, S2, 2-annotator sets	0.524	0.382	+0.142		
Follow-up (one change)	Variant	Qwen, 3 epochs	Δ	Seeds up	95% CI
12 epochs, single model, S2 [‡]	0.543 $\pm$ 0.015	0.462 $\pm$ 0.080	+0.082 [†]		
All of train, single model, S2	0.492 $\pm$ 0.045	0.476 $\pm$ 0.043	+0.017	6/10	
All of train, system, S2	0.575 (0.562)	0.543 (0.549)	+0.028		[−0.026, +0.092]

Table 3: The backbone swap (top) and two follow-ups (bottom) on dev. Single model: mean $\pm$ s.d. over seeds, Δ the paired mean difference. LA: logit adjustment. System: 10-seed ensemble + LA, τ by nested CV (in parentheses: mean over 10 CV splits); CIs are bootstraps over dev items. 2-annotator sets: the test regime. [‡]Seeds 0–2. [†]Reported only; the 12-epoch schedule was adopted on the train slice.




Figure 3: Dev S2 multi-reference macro-F1 of the DeBERTa-v3-large baseline, DeBERTa with the full question and 16 epochs, Qwen3-8B-Base with LoRA on the same input (3 epochs), and Qwen trained on all of train. (a) Single models on 10 paired seeds: open circles joined by thin lines are individual seeds; filled squares are means with ± 1 standard deviation. (b) The corresponding 10-seed systems (ensemble + logit adjustment, τ by nested CV; bar height is the mean over 10 CV splits), with published dev results as horizontal lines: TeleAI’s multi-call pipeline (0.617) and fine-tuned Qwen2.5-7B (0.495) (Shi et al., 2026), ChulaNLP’s DeBERTa top-5 + Kimi-K2 pipeline (0.52) (Peerachaidecho and Rutherford, 2026), and a human annotator scored against the other two (0.684).

rest), and the shorter half of dev scores 0.400 on S2 against 0.283 for the longer half.

A typical error shows the knowledge an answer can depend on (Figure 4): asked whether he would campaign against Senator Lieberman, the speaker replied “*I’m going to stay out of Connecticut.*” Recognising a near-answer requires knowing that Lieberman is a senator from Connecticut (our inference); errors of this kind motivated the knowledge hypothesis. Qwen ranks Dodging first, but logit adjustment then picks the rarer Declining to answer: the prior correction that raises macro-F1 over the set can cost an individual item.

Where the 8B classifier gains. Table 4 refutes our registered prediction. Per-class F1 rises in 7

of 9 classes, most in Claims ignorance (+0.403), Declining (+0.220), Dodging (+0.103) and Deflection (+0.089), while the commitment classes move least (General +0.059, Explicit +0.057, Implicit +0.040). Partial/half-answer is never predicted by either model. The in-set rate rises on unanimous items (+0.076) as much as on contested ones (+0.105, +0.048). The largest gains are on the Non-Reply classes, which turn on recognising what kind of statement an answer is; General remains the weakest frequent class (F1 0.322).

Ranking and the decision rule. Both models rank better than they decide. For Qwen (seeds 0–2), an acceptable label is the top prediction for 60.4% of items and in the top three for 93.1%, and the

Sub-question: Would you campaign against Senator Joe Lieberman on Iraq?

Full question: ... would you campaign against Senator Joe Lieberman, whose Republican candidate may support you, but he supports you too, on Iraq?

Answer: I'm going to stay out of Connecticut.

Annotators: Dodging, Dodging, Implicit (S1: Ambivalent)

	DeBERTa system	Qwen system
$\bar{p}$, top 3	General .373 Implicit .282 Deflection .205	Dodging .327 Implicit .189 Declining .170
Ensemble argmax	General $\times$	Dodging $\checkmark$
τ (item's fold)	0.70	0.85
S2 after LA	General $\times$	Declining $\times$
S1	Ambivalent $\checkmark$	Clear Non-Reply $\times$

Figure 4: One dev item through both 10-seed systems ($\checkmark$: in the reference set, or for S1 the gold level). The item was fixed in advance as the error discussed in §6. Qwen ranks an acceptable label first, but logit adjustment divides by the train prior (0.205 for Dodging, 0.042 for Declining to answer) and moves the decision to the rarer class.

in-set rate rises from 0.40 to 0.87 across fifths of dev from least to most confident; most remaining errors are a choice among a few candidates the model already ranks highly. Raw Qwen is skewed towards frequent classes: on seeds 0–2 it predicted General for only 24 items, although General is in 113 reference sets. Logit adjustment corrects this, adding +0.062 per Qwen model (0.476 to 0.538) against +0.005 for DeBERTa. A Qwen + DeBERTa probability mix, fitted inside the nested CV, adds nothing (0.543 against 0.549 for Qwen alone).

7 Conclusion and Plan to the Final Evaluation

A carefully controlled encoder for CLARITY stops at 0.384 Subtask 2 macro-F1 per model (0.405 as a system), and the evidence places the limit in the model rather than in the labels or the decision rule. Changing only the backbone to a LoRA-tuned 8B decoder improves both subtasks on all ten paired seeds and lifts the system to 0.543 / 0.746, still below the best multi-call pipeline (0.617) and the human reference (0.684). The gain is broad rather than concentrated where we predicted, which suggests that a larger pretrained model reads better what kind of response an answer is.

Plan. Table 5 gives the remaining work up to the final submission (31 October 2026); none of it has been run. It tests the remaining explanation,

	DeBERTa	Qwen	Δ
<i>per-class F1</i>			
Explicit	0.672	0.729	+0.057
Implicit	0.403	0.443	+0.040
Dodging	0.484	0.587	+0.103
General	0.263	0.322	+0.059
Deflection	0.287	0.377	+0.089
Partial/half-answer	0.000	0.000	+0.000
Declining to answer	0.381	0.601	+0.220
Claims ignorance	0.296	0.700	+0.403
Clarification	0.667	0.524	−0.144
<i>in-set rate by distinct annotator labels</i>			
one ($n = 125$)	0.529	0.605	+0.076
two ($n = 150$)	0.493	0.598	+0.105
three ($n = 33$)	0.688	0.736	+0.048

Table 4: Where the backbone gain falls: single models, mean over the same 10 seeds, dev S2. The registered prediction (a gain concentrated on the commitment boundary and on agreed items) is refuted.

Week	Planned work
Oct 5–11	Twelve-epoch Qwen on seeds 3–9 (a ten-seed system); all of train with twelve epochs as a one-change test.
Oct 12–18	Cascade: Qwen's top- k candidates, an open LLM choosing with definitions and boundary examples; prompts designed on the train slice.
Oct 19–25	Sweep the deferred fraction for an accuracy-vs-LLM-calls curve; cost per item.
Oct 26–31	Freeze systems; final ten-seed and two-annotator numbers; report and code release.

Table 5: Timeline to the final evaluation (all planned).

label meaning, with a cascade that sends only uncertain items, with the classifier's top candidates, to a larger LLM given label definitions and boundary examples, under the same protocol as above.

Limitations

Development-set evaluation. All results are on the 308-item development set, because the test labels were never released and the evaluation server has closed. The development set has three annotators per item and the test set two; we re-score with every two-annotator subset to approximate the test regime, but cannot measure the shift in answer length between the splits. With 308 items, system-level differences carry wide intervals: the gain of the 8B system has a 95% bootstrap interval of [0.054, 0.224], and the all-data variant could not be separated from its control.

Seeds and selection. Seed variance is large (the standard deviation of single-model Subtask 2 scores is 0.043 for the 8B classifier), so claims rest

on ten paired seeds; the twelve-epoch schedule was screened on three seeds only and is adopted, not established. Epochs are chosen on a 345-row slice of the training data, which is itself noisy. The one decision-rule parameter is fitted by nested cross-validation on the development set, which keeps the score held out but means the system uses development labels to fit one scalar.

Scope of the comparison. The backbone change also changes the training recipe that a backbone requires (learning rate, epochs, parameter-efficient fine-tuning), so “knowledge” is shorthand for everything that comes with a larger pretrained model. We cannot rule out that the pretraining data of the 8B model contains the public interview transcripts the dataset is drawn from. Published comparison points come from other teams’ papers and protocols and are not directly controlled.

Coverage. The rarest class (*Partial/half-answer*, 14 development reference sets) is never predicted by any of our models, and the study covers one dataset in one language (English, U.S. presidential interviews).

Ethical Considerations

The models classify the *form* of an answer relative to a question; they do not judge honesty, intent or the truth of what is said, and a prediction such as *Dodging* should not be presented as a verdict on a speaker. Annotators themselves disagree on many items (§6), and automated labels of political speech can be misused to support one-sided narratives; any deployment should report uncertainty and keep humans in the loop. The data are public interview transcripts released for the shared task. GPU experiments for the 8B classifier ran on rented single-GPU machines (up to four at once, for about two hours of wall-clock time); Appendix A reports their cost.

Acknowledgments

We thank the organisers of SemEval-2026 Task 6 for the data and the evaluation tools.

References

Pengcheng He, Jianfeng Gao, and Weizhu Chen. 2023. DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing. In *The Eleventh*

International Conference on Learning Representations (ICLR).

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. LoRA: Low-rank adaptation of large language models. In *The Tenth International Conference on Learning Representations (ICLR)*.

Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. 2017. Focal loss for dense object detection. In *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pages 2980–2988.

Aditya Krishna Menon, Sadeep Jayasumana, Ankit Singh Rawat, Himanshu Jain, Andreas Veit, and Sanjiv Kumar. 2021. Long-tail learning via logit adjustment. In *The Ninth International Conference on Learning Representations (ICLR)*.

Siva Nuthakki, Sanjay Pulagam, and Sai Woon. 2026. Team evaluators at SemEval-2026 task 6: Instruction-tuned LLMs for clarity and evasion classification in political interviews. In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 3442–3445, San Diego, California, USA. Association for Computational Linguistics.

Wisarat Peerachaidecho and Attapol T. Rutherford. 2026. ChulaNLP at SemEval-2026 Task 6: A hybrid BERT–LLM framework for political response clarity and evasion detection. In *Proceedings of the 20th International Workshop on Semantic Evaluation (SemEval-2026)*, pages 1874–1881. Association for Computational Linguistics.

Qwen Team. 2025. *Qwen3 technical report*. Preprint, arXiv:2505.09388.

Jiawei Ren, Cunjun Yu, Shunan Sheng, Xiao Ma, Haiyu Zhao, Shuai Yi, and Hongsheng Li. 2020. Balanced meta-softmax for long-tailed visual recognition. In *Advances in Neural Information Processing Systems 33 (NeurIPS)*.

Lingling Shi, Haoyu Jin, Shiquan Wang, Ruiyu Fang, Shuangyong Song, and Xuelong Li. 2026. TeleAI at SemEval-2026 Task 6: A confidence-aware multi-stage reasoning framework with chain-of-thought. In *Proceedings of the 20th International Workshop on Semantic Evaluation (SemEval-2026)*, pages 1591–1609. Association for Computational Linguistics.

Taleef Tamsal and Jonathan Rusert. 2026. PFW at SemEval-2026 task 6: Multi-seed DeBERTa ensembles for political response clarity and evasion classification. In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 1518–1525, San Diego, California, USA. Association for Computational Linguistics.

Amruth Tetakali and Lavanya Tetakali. 2026. SilkPeak at SemEval-2026 task 6: When politicians dodge — unmasking evasion in political interviews through

[joint multi-task transformer learning](#). In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 2020–2025, San Diego, California, USA. Association for Computational Linguistics.

Konstantinos Thomas, Giorgos Filandrianos, Maria Lymperaïou, Chrysoula Zerva, and Giorgos Stamou. 2024. “I never said that”: A dataset, taxonomy and baselines on response clarity classification. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 5204–5233, Miami, Florida, USA. Association for Computational Linguistics.

Konstantinos Thomas, Giorgos Filandrianos, Maria Lymperaïou, Chrysoula Zerva, and Giorgos Stamou. 2026. [SemEval-2026 Task 6: CLARITY – unmasking political question evasions](#). *Preprint*, arXiv:2603.14027.

Tai Tran Tan and An Dinh. 2026. [ttda704 at SemEval-2026 task 6: Structured chain-of-thought prompting for political evasion detection](#). In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 2817–2829, San Diego, California, USA. Association for Computational Linguistics.

Nawar Turk, Lucas Miquet-Westphal, and Leila Koseim. 2026. [CLaC at SemEval-2026 task 6: Response clarity detection in political discourse](#). In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 1890–1902, San Diego, California, USA. Association for Computational Linguistics.

Alexandra N. Uma, Tommaso Fornaciari, Dirk Hovy, Silviu Paun, Barbara Plank, and Massimo Poesio. 2021. Learning from disagreement: A survey. *Journal of Artificial Intelligence Research*, 72:1385–1470.

A Implementation Details

Encoder. Table 6 lists the DeBERTa settings; baseline and final models differ only in the input and the number of epochs. One detail mattered: transformers 5.x loads a checkpoint in its stored dtype, and DeBERTa-v3’s published weights are fp16, in which its attention overflows; the model then trained smoothly to the label prior and predicted Explicit for every item, so we load the weights in fp32 explicitly.

Setting	Value
Backbone	microsoft/deberta-v3-large, all weights trained
Head	[CLS] $\rightarrow$ dense + tanh $\rightarrow$ dropout 0.1 $\rightarrow$ linear (9)
Input, baseline	sub-question (≤ 96 tokens) [SEP] answer; 512 tokens
Input, final	“Sub-question: ... Full question: ...” (≤ 256 tokens) [SEP] answer; 1,024 tokens
Truncation	answer cut from the right
Loss	cross-entropy, no smoothing, no class weights
Optimiser	AdamW, $\beta = (0.9, 0.98)$, $\epsilon = 10^{-6}$, weight decay 0.01 (none on biases and LayerNorm)
Learning rate	$10^{-5} \times 0.95^k$ for a layer k levels below the head; head and pooler 10^{-4}
Schedule	cosine, warm-up 10% of steps
Batch	16, no accumulation; gradient-norm clipping 1.0
Epochs	8 (baseline) or 16
Precision	fp32 weights, bf16 autocast; gradient checkpointing
Epoch selection	best macro-F1 on the 345-row train slice; last epoch for all-of-train runs
Seeds	0–9 (all of train: 0–4)
Logit adjustment	$\tau \in [-1, 2]$, 61-point grid; prior = train label frequencies; 5-fold nested CV on dev
Hardware	one shared RTX PRO 6000 (96 GB); peak memory about 9 GB at 512 tokens, 15–17 GB at 1,024

Table 6: DeBERTa-v3-large training and decision settings.

LLM classifier. Table 7 lists the Qwen3-8B-Base settings; everything not listed matches the final encoder configuration. Each run used one NVIDIA RTX PRO 6000 (96 GB) on a rented virtual machine. A 3-epoch seed took 15–21 min and a 12-epoch seed about 59 min; the ten-seed round cost INR 1,518.79 of GPU time.

Code and records. Code, experiment definitions, the experiment log with every pre-registered hypothesis, and the raw analysis outputs behind every number are at <https://github.com/IamSid44/>

Setting	Value
Backbone	Qwen3-8B-Base, bf16
Adapter	LoRA, $r = 16$, $\alpha = 32$, dropout 0.05
LoRA targets	q, k, v, o, gate, up, down_proj
Head	new 9-way linear, last non-pad token, trained in full
Precision	fp32 trainable weights, bf16 autocast
Optimizer	AdamW, lr 10^{-4} , weight decay 0
Gradient clipping	norm 1.0
Schedule	cosine, linear warmup 5% of steps
Epochs	3 (12 in the follow-up)
Batch	4×4 accumulation = 16
Steps per epoch	194 (3,103 training rows)
Loss	cross-entropy, unweighted
Input	question part ≤ 256 tokens; total $\leq 1,024$ incl. closing question
Batching	length-sorted within chunks of 50 batches; order fixed by (seed, epoch)
Epoch selection	best macro-F1 on the 345-row slice (all-of-train runs: last epoch)
Padding check	abort if relative logit difference > 0.2
Seeds	0–9 (12 epochs: 0–2)

Table 7: Hyperparameters of the LLM classifier.

CLARITY-Unmasking-Political-Question-Evasions. Per-seed probabilities, metrics and LoRA adapters are at <https://huggingface.co/siddharthg44/clarity-semeval26>, and training curves at <https://wandb.ai/iamsid44-iiit-hyderabad/clarity-semeval26>. Software: Python 3.12, PyTorch 2.11, Transformers 5.17, PEFT 0.20.