Evasion Is a Relation: A Confidence-Routed Cascade for Political Question Evasion, and What It Measures

Vidvathama R 2024122002 **Siddarth Gottumukkula** 2023102040 **Sanjana Reddy Vonteri** 2026901007 **Shashikanta Sahoo** 2026900007

Team Nier_ANLP
International Institute of Information Technology, Hyderabad

Abstract

SemEval-2026 Task 6 (CLARITY) asks how a politician's answer responds to one interview question, as one of nine evasion types (Subtask 2) or three clarity levels (Subtask 1). We ask five questions in turn, each registered before its experiments and each raised by the answer to the one before. A controlled encoder study finds that architectural changes do not help, while swapping only the backbone for a LoRA-tuned Qwen3-8B does (+0.092 Subtask-2 macro-F1 per model, on ten of ten paired seeds). That classifier ranks better than it decides, so a 27B LLM chooses among its top four labels for the items it is least sure of (30% of a training slice, 36% of dev). The resulting cascade scores 0.650 (Subtask 2) and 0.788 (Subtask 1) on the development set with 1.4 LLM calls and about 1.6k LLM input tokens per item, against 0.617 and 0.812 at 1.94 calls and about 7.3k tokens for the winning system. A distilled single-model variant matches the ten-model classifier at a quarter of its GPU time. We then freeze the cascade and use it as an instrument. The labels of one answer change with its sub-question far more often than annotator noise explains, and pairing an answer with an unrelated question cuts the Clear-Reply share from 25% to 1%, so evasion is a relation between a question and an answer. When two open LLMs answer the Presidents' questions as a neutral analyst, their evasions show no clear link to the Presidents' (odds ratio 1.18), but in the President's role the models answer directly the questions the Presidents answered directly (2.74, exploratory), and a follow-up turns 66% of their evasions into replies. Finally, running the classifier in FP8 costs it 0.045 Subtask-2 macro-F1 with all parameters frozen, while the full cascade loses 0.003.

1 Introduction

Politicians often answer the question they would have liked to be asked. Thomas et al. (2024) turned this observation into a two-level taxonomy of response clarity and evasion, and SemEval-2026 Task 6 (Thomas et al., 2026) made it a shared task. The task was won by pipelines that call a large language model (LLM) several times per item (Shi et al., 2026), while fine-tuned encoders levelled off near 0.50 macro-F1 on the nine-way subtask. A leaderboard shows which system is best, but not why the others stop where they do, nor what a good classifier of evasion could be used to find out.

This paper addresses those two questions in five stages (Figure 1). Figure 2 shows the system it builds and how it is used. At each stage we wrote down a hypothesis, a prediction and a decision rule before running anything, and the result of each stage set the question of the next.

Our contributions follow the same order. First, controlled evidence that a stronger pretrained backbone helps where structural changes do not. Second, a cascade at 0.650 Subtask-2 macro-F1 with 1.4 LLM calls per item, and a distilled variant at a quarter of the classifier's GPU time. Third, evidence that evasion is relational, and an exploratory LLM audit in which, in the President's role, models answer directly what the Presidents answered directly. Fourth, a measurement of what reduced precision costs the cascade. Every system comparison is on the 308-item development set, since the test labels were never released. Many registered hypotheses failed. We report them where they fell, because each failure changed what we did next.

2 Task, Data and Protocol

Task. Each item is a triple $x = (q, Q, a)$ from a U.S. presidential interview (Thomas et al., 2024): a sub-question q (median 15 tokens), the journalist's full turn Q it was split from (62 tokens) and the President's whole answer a (266 tokens). Subtask 2 (S2) asks for one of nine evasion types $\mathcal{Y}_9$, describing how a responds to *that* sub-question, and each type fixes one of three clarity levels for

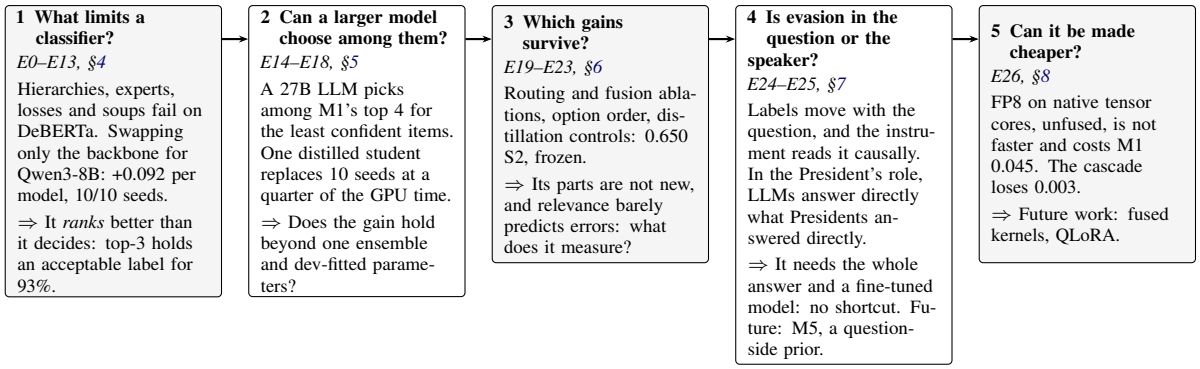


Figure 1: The research path. Each stage answers one question, and its answer (last line of each box) raises the next. Experiment ids refer to the ledger in Appendix A.

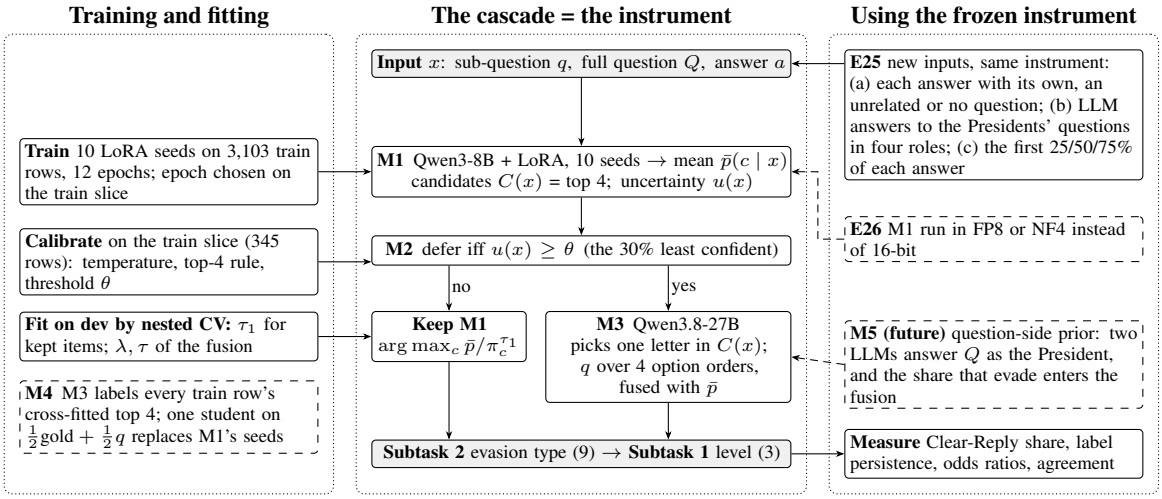


Figure 2: The whole system. Left: what is trained or fitted, and on which data. Middle: the cascade, frozen after E23 as the *instrument*. Right: its use, on new inputs (E25) or with a lower-precision M1 (E26), always with the same settings. Dashed boxes are optional (M4, E26) or proposed future work (M5).

Evasion type (S2)	Clarity level (S1)	Train %	Dev sets
Explicit	Clear Reply	30.5	115
Implicit	Ambivalent	14.2	99
Dodging	Ambivalent	20.5	96
General	Ambivalent	11.2	113
Deflection	Ambivalent	11.0	51
Partial/half-answer	Ambivalent	2.3	14
Declining to answer	Clear Non-Reply	4.2	17
Claims ignorance	Clear Non-Reply	3.5	15
Clarification	Clear Non-Reply	2.7	4

Table 1: The taxonomy. Train % is the share of the 3,448 training labels. Dev sets is the number of the 308 dev items whose reference set contains the type.

Subtask 1 (S1, Table 1). Like the top systems, we predict the nine-way type and map it to the clarity level.

Data. Train has 3,448 items with one label each. Dev has 308 items with three annotators: 125 are unanimous, 150 split two to one and 33 receive three labels, so the reference set G_i of an item holds one to three types. The 237 test items have two annotators, but their labels were never released, so every comparison is on dev, and we also rescore dev against each of its three two-annotator subsets (the test regime). One property shapes the whole paper: 69% of training rows share their answer with another sub-question of the same turn (*siblings*).

Metric. S2 is scored by multi-reference macro-F1: a prediction is a true positive if it lies in G_i, and every class counts equally. It rewards both the *in-set rate* and *coverage* of rare classes. Always predicting Explicit scores 0.060, TF-IDF with logistic regression 0.257, and a human annotator against the other two 0.684.

Protocol. Epochs, calibration and routing thresholds are chosen on the *train slice*, a fixed stratified 10% of train (345 rows). The few parameters that

must be fitted on dev are chosen by 5-fold nested cross-validation, and we report the first split and the mean over ten splits. A *system* averages seeds and applies logit adjustment (Menon et al., 2021), $\hat{y} = \arg\max_c \bar{p}(c \mid x)/\pi_c^\tau$ with π_c the train prior and τ the only fitted parameter. Claims rest on ten paired seeds, and a change must first pass a three-seed screen (+0.015, at least two of three seeds up). Before each experiment we registered its hypotheses, predictions and decision rules. Later changes were logged as amendments before results were seen. From §7 on, every label comes from one *frozen* system, so no parameter changes between experiments.

3 Related Work

The task and its systems. Thomas et al. (2024) introduced the clarity and evasion taxonomy and found that predicting the nine-way type and mapping it up beats predicting clarity directly. In the shared task (Thomas et al., 2026), fine-tuned encoders levelled off near 0.50 on test S2 and the strongest systems prompted LLMs. TeleAI (Shi et al., 2026) won with three confidence-routed chain-of-thought stages on DeepSeek-V3 (1.94 calls per item, dev 0.617 S2 and 0.812 S1), and also reported a fine-tuned Qwen2.5-7B at 0.495. ChulaNLP (Peerachaidecho and Rutherford, 2026) passed a DeBERTa model's top five labels to a prompted LLM (dev 0.52). The parts of our cascade are not new. We add a controlled account of which parts matter, and its use as an instrument.

Cascades and LLM choosers. Sending only some inputs to a larger model, as FrugalGPT (Chen et al., 2023) does for LLM APIs, is the idea behind our router. LLMs choosing among lettered options are sensitive to option order (Zheng et al., 2024), which we measure and correct, and we distil the adjudicator into the classifier with soft targets (Hinton et al., 2015).

Evasion as a property of the exchange. Conversation analysis defines evasion relative to the question: an answer evades by failing to meet the agenda the question sets (Clayman, 2001). Alvarez and Morrier (2025) score answer quality by how identifiable an answer is among others given its question. We test whether such relevance captures the taxonomy's labels, and we use LLMs as speakers with no political stake under the same instrument.

4 Stage 1: What Limits a Fine-Tuned Classifier?

Question. Label noise, a decision rule that does not fit the metric, missing knowledge and labels too subtle for 3,448 examples each predict a different fix. We tested these explanations on DeBERTa-v3-large (He et al., 2023), fully fine-tuned with a nine-way head, changing one thing at a time against a named control (Figure 3a).

Result: structure fails, the decision rule and training length help. Every system averages the softmax outputs of ten seeds that differ only in their random seed, which reduces variance, and then applies a one-parameter logit adjustment, $\hat{y} = \arg\max_c \bar{p}(c \mid x)/\pi_c^\tau$. Dividing by the train prior π_c lifts the rare classes that an argmax underpredicts and that macro-F1 weights equally. It raised an encoder system from 0.365 to 0.438 (E4), whereas nine free class weights reached only 0.394 and overfit dev (in-fold gain +0.098). Longer training and the full journalist question improved the model itself: 0.315 to 0.384 per model, 9 of 10 seeds up. Every structural idea failed: hierarchies did not beat a flat softmax that already scores all nine labels jointly, rebalanced losses over-corrected towards rare classes (Explicit F1 0.68 $\rightarrow$ 0.32), and boundary experts rarely overturned decisions the flat model made from the same input. The clearest case is the definition re-ranker: gold was in the encoder's top three for 90% of items, yet a second encoder given label definitions chose worse.

Why: the model, not the labels. Of the encoder ensemble's predictions, 44.5% fell outside the reference set, 53 of them on items where all three annotators agreed, while a human annotator scores 0.684 against the other two. So label noise does not explain the gap, and nothing trained on the same 3,448 rows moved the system.

Hypothesis E13: the gap is knowledge. We replaced only the backbone with Qwen3-8B-Base (Yang et al., 2025), adapted with LoRA (Hu et al., 2022) (rank 16, all linear layers) and a new linear head on the last token, keeping the data, input, loss and selection rule. We expected the gain to concentrate on the commitment boundary (Explicit–Implicit–General), where the encoder's errors sit, and on unanimous items. The gain was +0.092 S2 and +0.095 S1 per model on all ten seeds, more than the whole encoder track gained. The 10-seed system reached 0.543 S2 and 0.746 S1 (95% CI of the gain over DeBERTa [+0.054, +0.224]). The

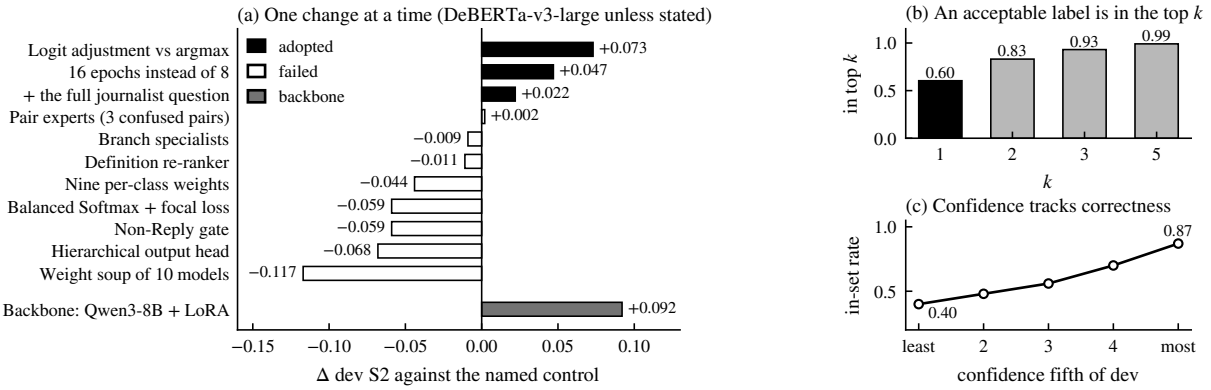


Figure 3: Stage 1 on dev S2. (a) Each change against its named control (E0–E13): per model for training and input changes, per system for the rest. (b, c) The Qwen3-8B classifier ranks better than it decides, and correctness rises with confidence (3 epochs, seeds 0–2).

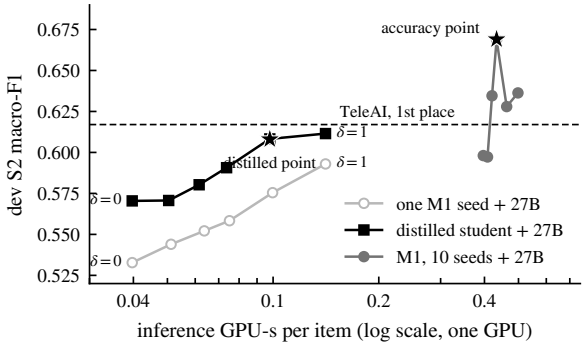


Figure 4: The cascade on dev: S2 against inference GPU-s per item, for deferral rates $\delta = 0, 0.1, 0.2, 0.3, 0.5, 1$ (left to right on each curve, first scoring, one option order). All times are on one GPU (Appendix B). TeleAI's GPU time is not published.

location was not as expected: the gain was broad (seven of nine classes), largest on Claims ignorance (+0.40) and Declining (+0.22), the classes that depend on recognising *what kind of statement* an answer is.

What it raised. The 8B classifier ranks far better than it decides, and it knows when it is unsure (Figure 3b–c). Most remaining errors are choices among a few candidates it already ranks highly, made where it is unsure.

5 Stage 2: Can a Larger Model Choose Among the Candidates?

Question. If M1's errors are choices among a few top candidates made where it is unsure, can a larger model consulted only there make better choices?

M1, the candidate generator (E14). On ten seeds the 12-epoch classifier scores 0.525 ± 0.042 per model (+0.049 over 3 epochs, 8/10 up) and 0.598 S2 / 0.743 S1 as a system (M1, Figure 2). Three registered checks shaped the next module. An oracle re-deciding the 30% least confident items within the top three would reach 0.727, but a *random* choice there lowers S2 to 0.587: M1's first choice carries information, so an adjudicator must be *fused* with M1, not substituted for it. Candidate sets defined by probability mass did not transfer from the slice to cross-fitted rows (coverage 0.98 vs 0.76) while rank-based sets did, so $C(x)$ is the top four. And five seeds did not score within 0.01 of ten as predicted (0.555 vs 0.598), so M1 kept ten seeds and became the dominant cost.

M3, the adjudicator (E15). Qwen3.8-27B-FP8, run locally without thinking, reads the label definitions, a guide to the most-confused pairs and the item, and answers with one letter among the four candidates, and $q(c \mid x)$ is the softmax over the letters' log-probabilities. On deferred items the label is $\arg\max_{c \in C(x)} \lambda \log \bar{p} + (1 - \lambda) \log q - \tau \log \pi_c$, with (λ, τ) fitted by nested CV. Other items keep M1's answer. A screen on 300 cross-fitted training rows refuted two registered design assumptions: retrieved examples *hurt* (fused accuracy 0.353 with none, 0.330 with ten), and a 304B mixture-of-experts model was worse than the local 27B (0.317).

Result, and a correction. The first dev analysis showed no gain. Explaining it, we found a bug that shared one τ between deferred and kept items and so switched off M1's logit adjustment on the 64% of items meant to keep M1's answer. Corrected, deferring the 30% least confident items (by the slice threshold, 36% of dev) gives 0.669 S2 (0.660 over ten splits) and 0.798 S1, +0.067 [+0.009, +0.124] over M1 with the fitted param-

eters held fixed (Figure 4), with 33 items moving into the reference set and 9 out (McNemar $p = 0.0003$). Deferring *every* item is worse (0.609 over ten splits): where M1 is confident, the LLM is the weaker reader.

Where the cost goes (E16). Timed on one GPU, an M1 seed costs 0.040 GPU-s per item and a 27B call 0.101, so the ten seeds are 92% of the cascade's inference time and the 27B only 8%.

M4, distillation for cost (E17–E18). M3 relabelled every training row over its cross-fitted top four, and one student learned from $\frac{1}{2}$gold $+\frac{1}{2}q$, a weight fixed in advance (Hinton et al., 2015). The student beats M1 per model (0.573 vs 0.525, 9/10 seeds up) and is far better calibrated (mean top probability 0.58 vs 0.84), yet it ranks its errors no better (AUROC 0.686 for both), against our prediction. With M3 it missed the registered target at $\delta = 0.3$ (0.591 vs 0.614). At $\delta = 0.5$ and with one option order it matches the ten-seed M1 alone (0.599 against 0.598 S2, Table 2) at 0.098 GPU-s per item, a quarter of M1's time. Ensembles of students add almost nothing (+0.005 from one to ten, vs +0.065 for M1 seeds): students trained towards the same targets lose the diversity seed ensembles feed on. We froze an *accuracy* point (M1 $\times$ 10 + M3, $\delta = 0.3$) and a *distilled* point (one student + M3, $\delta = 0.5$).

What it raised. The accuracy point rested on one ensemble (smaller ensembles gained less from M3), on parameters fitted on 308 items, and, for the distilled point, on a δ chosen after the registered one failed.

6 Stage 3: Which Gains Survive?

Question. The Stage 2 gain rests on one ensemble, one adjudicator and one prompt. Each simpler alternative predicts a specific failure, so we registered one test per alternative (Figure 5), decided on training rows or held-out folds, never on dev.

Result. Three findings change how the cascade should be read. On 3,448 cross-fitted training rows the cascade gains a small but significant +0.014 [+0.002, +0.026], less than on dev, where its parameters were fitted. The gain is carried by the two parts a simpler system lacks: confidence routing, which no random routing matches (Figure 5a), and fusion (the 27B alone over all nine labels scores 0.591, below M1's 0.598, and deciding the deferred items alone scores 0.631). The 27B breaks as many confident decisions as it fixes and helps only where M1 is unsure (b), which is why a learned fix/break router lost to confidence at all four budgets. And the 27B's decisions depend on more than the item. LLMs choosing among lettered options are known to prefer some positions over others (Zheng et al., 2024), and the 27B prioritises option A: in taxonomy order it picks A for 40% of items, against 22–29% in the other orders and 25% by chance (Figure 5c). Because slot A usually holds Explicit in taxonomy order, a preference for A is also a preference for Explicit. Two orders of the same four candidates give the same pick on only 72% of dev items, so we average four, which costs four calls and moved the system by +0.011 at split 0 but +0.002 over ten splits. Two results also changed what we claim. A student taught by M1's own soft targets came within the registered margin of the M3-taught one, so the distilled point is a soft-target effect, not "the LLM's knowledge". Transcript artifacts in 97% of test answers changed 15.6% of M1's labels when injected into dev, so test predictions use normalised text.

The frozen instrument (E23). Re-scoring every record on one serving setup of the 27B with order-averaged q gives the final operating points (Table 2), slightly lower than the Stage 2 scores because all 27B decisions now come from one scoring run. Averaging four orders costs four 27B calls per deferred item, so the accuracy point, 0.650 S2 (0.639 over ten splits) and 0.788 S1, makes 1.43 calls per item and takes 1.36$\times$ M1's GPU time (Table 2 gives the one-order variants). Against the winning system it is higher on S2 and lower on S1, with three quarters of its calls and a fifth of its LLM input tokens (its GPU time is not published). Every setting of the accuracy point is fixed in one file (Appendix B). From here on it is the *instrument*, and nothing in it is refitted.

What it raised. The cascade's gain is now explained, but not what its labels mean: question–answer relevance, tried as a router feature (E22), barely predicted its errors ($\rho = +0.095$). If the label is not a property of the answer's relevance, what is it a property of?

7 Stage 4: What Does the Instrument Measure?

7.1 Labels across the sub-questions of one answer (E24)

Question. If evasion were a property of an answer's style, one answer would get one label what-

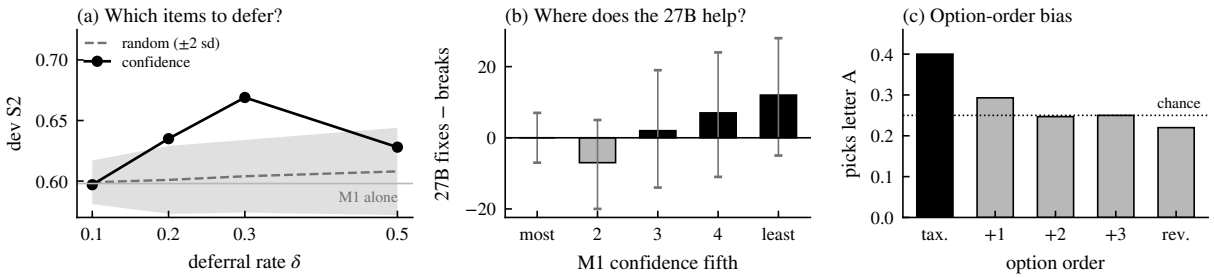


Figure 5: Stage 3 (E19–E22), registered tests of the Stage 2 result. (a) Confidence routing against 200 random routings. (b) Net fixes of the 27B by M1's confidence, on 1,067 leak-free training rows. (c) How often the 27B picks the first option, by order.

System (dev)	S2	S1	calls	GPU-s
DeBERTa-v3-large, 10 seeds	0.405	0.648	0	–
M1: Qwen3-8B, 10 seeds	0.598	0.743	0	0.397
Accuracy point (M1 + M3)	**0.650**	**0.788**	1.43	0.538
one option order	0.639	0.787	0.36	0.433
Distilled point (student + M3)	**0.608**	**0.785**	2.30	0.266
one option order	0.599	0.784	0.57	0.098
TeleAI, 1st place	0.617	0.812	1.94	n/a
Human vs. two others (2-ann.)	0.684	–	–	–

Table 2: Final systems (E23, one serving setup, nested-CV split 0). Calls: 27B calls per item (four per deferred item with order-averaged q). GPU-s: inference per item on one GPU (Appendix B). TeleAI uses ∼7.3k LLM input tokens per item, the accuracy point ∼1.6k.

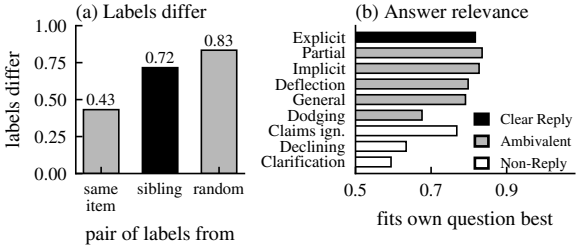


Figure 6: Evasion is a relation (E24). (a) Share of dev label pairs from two annotators that differ: same item, two sub-questions of one answer, random items. (b) Share of three random questions under which the answer is less likely than under its own, by gold label (train): Partial, Implicit, Deflection and General answers are as relevant as Explicit ones.

ever it was asked. Answers shared by several sub-questions test this (CPU only, existing outputs).

Result. Of the 3,079 training pairs of siblings (two sub-questions, one answer), 61.5% carry different evasion types and 40.6% different clarity levels. This is not annotator noise: on dev, moving to a sibling's sub-question changes the label almost as often as moving to an unrelated exchange, and far more often than two annotators disagree on the same item (Figure 6a). The label belongs to the pair, not to the answer. M1 already behaves accordingly, repeating its label on only 19% of the 58 dev sibling pairs with disjoint annotator sets.

Relevance is not answerhood. Political scientists score answer quality by how specific an answer is to its question (Alvarez and Morrier, 2025). We computed this with the frozen Qwen3-8B as the gain in the answer's log-likelihood under its own sub-question over three random ones. It is ordered across the nine labels (Kruskal–Wallis $p = 2 \times 10^{-12}$) but barely discriminative (AUROC 0.566 for Clear Reply against the rest): Partial, Implicit, Deflection and General answers fit their own question as well as Explicit ones do (Figure 6b). An answer can be *about* the question without *answering* it, which is the distinction the taxonomy draws.

7.2 Does the instrument read the question causally? (E25a)

Registered prediction (H25a-1). An answer paired with a question from another interview has, by the taxonomy's definition, not replied to it. If the instrument models the relation, the Clear-Reply share should fall by at least half and at most 50% of items should keep their label. If it labels answer style, labels should persist. This is not obvious: training never showed a mismatched pair.

Result. Held (Figure 9a, appendix). With an unrelated question the Clear-Reply share falls from 0.253 to 0.013, only 18.8% of items keep their label, and 62 of the 78 Clear Replies become Dodging. The label-free replication on test runs in the same direction (0.338 to 0.110). Without any question the share *rises* to 0.640: without a question, nothing is evaded. The instrument models the relation, the precondition for using it on new speakers.

	Registered prediction	Result	Verdict
a-1	unrelated question: CR $\leq \frac{1}{2}$	$0.253 \rightarrow 0.013$	held
b-1	no stake: OR ≥ 1.5	1.18	failed
b-2	stake: P $-$ N $\geq +0.10$	-0.031	failed
b-3	pursuit: $\geq$50% replies	66%	held
b-4	notes: indirect $-$ non $\geq +0.15$	+0.113	failed
c	half answer: $\geq$70% agree	66%	failed
d-1	frozen probe: S2 ≥ 0.45	0.397	failed
d-2	probe: siblings same $\geq$30%	45%	held

Table 3: E25's registered hypotheses (dev, frozen instrument). CR is the Clear-Reply share and OR the odds ratio, shown beside each pair in Figure 7b.

7.3 Does the question or the speaker drive evasion? (E25b)

Design. If evasion is relational, it could be driven by the question (some questions are hard to answer for anyone) or by the speaker (Presidents evade strategically). We put the 308 dev sub-questions (191 interview turns) to two open instruction-tuned LLMs, Mistral-7B-Instruct-v0.3 (Jiang et al., 2023) and Qwen3-8B with thinking off (Yang et al., 2025), in four conditions that each differ from P in one factor. **N** is a neutral analyst with no stake, **P** the President in that interview on its date, **F** P followed by "With respect, you haven't answered the question", and **B** P with the President's real answer as private notes. The frozen instrument labelled all 2,464 answers against their sub-questions. Statistics use a cluster bootstrap over turns and Mantel–Haenszel odds ratios pooled over models.

Results. Both registered hypotheses about the question and the stake failed (Table 3, Figure 7): a neutral analyst does not evade more on the questions the Presidents evaded (OR 1.18), and playing the President does not raise the evasion *rate* (-0.031). In the President's role, though, the same models' evasions line up with the Presidents' (OR 2.74 [1.87, 4.17], with 3.0 for Mistral and 2.5 for Qwen). We did not register this comparison, so it is exploratory. In plain rates (Figure 7b), the neutral analyst evades on 56% of the questions the Presidents evaded and on 52% of those they answered. In the President's role the models evade on 58% and 34%. The move therefore comes from the questions the Presidents answered directly, which the models now also answer directly, not from the questions the Presidents dodged. A neutral analyst, to whom the questions are not addressed, evades about half of them whatever the President did. The role changes *where* models evade, not how often (52% against 55%). A question's answerability shows only for a speaker who can answer it, and our data cannot tell this apart from strategic evasion. A follow-up turns 66% of evasive answers into clear replies. With the President's answer as notes, indirect answers become replies more often than non-answers, by less than the registered margin.

Is the instrument valid on LLM text? The registered human check was dropped for lack of time. Without labels, the instrument is *more* confident on LLM answers than on the Presidents' (seed agreement 0.81 vs 0.73), though confident errors remain possible.

7.4 Is there a shortcut? (E25c–d)

A shortcut would exist if evasion were visible early in an answer or encoded in the frozen base model. Neither holds (Figure 9b–c). Half an answer agrees with the full-answer label on only 66% of items: refusals show early, General and Dodging only at the end. Linear probes on the frozen Qwen3-8B-Base reach 0.397 S2 at best, against 0.525 for one fine-tuned seed, and they give disjoint siblings the same label 45% of the time, against 19% for M1. Fine-tuning is what ties the label to the question.

What it raised. The question thus carries evasion signal that the cascade never uses, which we leave as a question-side prior (M5 in Figure 2, Limitations). Measuring evasion also needs the whole answer and the fine-tuned ten-seed ensemble, so the remaining cost lever that leaves the instrument's labels unchanged is numeric precision.

8 Stage 5: Can the Instrument Be Made Cheaper?

Design and predictions. The ten seeds are three quarters of the instrument's GPU time. Unlike distillation (§5), reduced precision changes only the arithmetic, so we ran the frozen M1 at reduced precision and compared each format with 16-bit on the same GPU type. FP8 (8-bit weights and activations) was compared through the full instrument, NF4 (4-bit weights, as in QLoRA, Dettmers et al., 2023) at the level of M1. We predicted that FP8 would keep $\geq$98% label agreement and S2 within 0.01 (H26-1) and be at least 1.3$\times$ faster at the same batch size (H26-2). The non-obvious prediction, H26-4, was that FP8's flips would fall where the 27B re-decides, so the instrument would lose less than M1.

Result. Both FP8 predictions about M1 failed

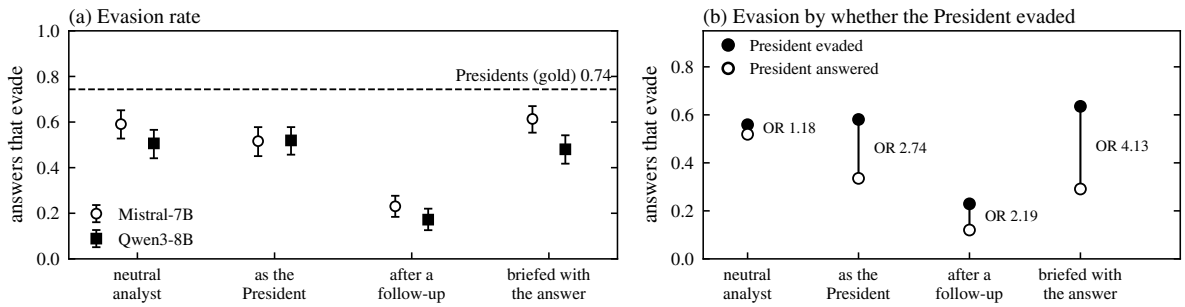
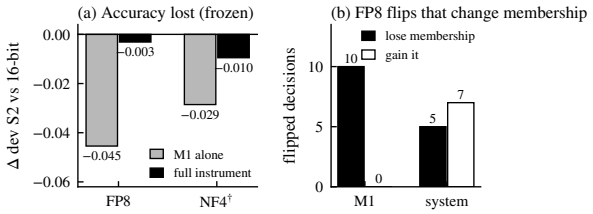


Figure 7: E25b, labelled by the frozen instrument. (a) Share of LLM answers that evade, with 95% cluster-bootstrap intervals (dashed line: the Presidents' gold rate). (b) The association as plain rates: the share of answers that evade where the President evaded (filled) and where the President answered (open), pooled over the two models, with the odds ratio beside each pair.



(c) M1, 16 bit → reduced	FP8	NF4†
Label agreement (10 seeds)	96.1%	86.7%
Throughput (items/s)	43.1 → 35.9 (**0.83**×)	7.7 → 6.3 (**0.82**×)
Batch-1 latency p50 (ms)	33.6 → 30.2 (0.90×)	169 → 259 (**1.53**×)
Peak memory (GB)	21.9 → 20.9 (**0.96**×)	16.8 → 11.3 (**0.67**×)
S2 lost, frozen parameters	0.045	0.029

Figure 8: E26 on dev, each format against 16 bit on the same GPU type. (a) S2 lost, both with frozen parameters. (b) FP8 flips that change reference-set membership. (c) Cost of M1 in reduced precision. FP8: one RTX PRO 6000, batch 32 for both. NF4: two RTX A4000, 16 bit split over two cards at batch 8 and NF4 on one card at batch 16. For scale, the weights of an 8B model alone take about 16, 8 and 4.5 GB in 16 bit, FP8 and NF4. †Lower bound for the full instrument: 27B records missing for 60 NF4 items.

(Figure 8c): FP8 changes 3.9% of the ensemble's argmax labels (17 of M1's final decisions) and lowers M1's S2 by 0.045 with the parameters frozen. We ran FP8 on the native FP8 tensor cores of an RTX PRO 6000 (Blackwell). As implemented it still lowers batched throughput to 0.83×, because activations are quantized in unfused operations around each FP8 matrix product. Reported FP8 speed-ups rely on kernels that fuse these steps into the matrix product, which we did not use. FP8 also barely changes peak memory (0.96×), since the embeddings, norms and head stay in 16 bit. Applied to the accuracy point's one-GPU cost, the slower M1 would raise it from 0.538 to 0.619 GPU-s per item (+15%). H26-4 held on its two registered counts (Figure 8a,b): 71% of FP8's 17 flips land on items the 27B re-decides, twice the chance rate, and the system flips 16 decisions against M1's 17. All 10 of M1's membership-changing flips lose reference-set membership, whereas the system's go both ways (5 lose, 7 gain), so with the same frozen parameters it loses 0.003 against M1's 0.045. The cascade does not undo the flips but turns a one-sided loss into noise. NF4 changes more (86.7% ensemble agreement) and meets our registered QLoRA trigger. It saves memory but not time: peak memory falls to 0.67× (16.8 to 11.3 GB), short of the 0.4× we predicted, while batch-1 latency rises from 169 to 259 ms, because the 4-bit weights are stored compactly but computed in bf16. Its throughput (0.82×) is not like for like, since the 16-bit reference ran split over two cards. Its cascade comparison is incomplete (27B records are missing for 60 of 136 deferred items), so we draw no conclusion from it. Precision is not a cost lever with these kernels.

9 Conclusion

A classifier on this task is limited by knowledge, not structure, and ranks better than it decides. A 27B adjudicator consulted only where it is unsure lifts dev S2 to 0.650 with fewer LLM calls and tokens than the winning system. As a frozen instrument, it shows that evasion is a relation between question and answer, and that in the President's role LLMs answer directly what the Presidents answered directly (exploratory), which motivates a question-side prior as the next module. Run in FP8 without fused kernels, the classifier loses 0.045 S2 and is not faster, while the cascade loses 0.003.

Ethical Considerations

The instrument labels the form of an answer relative to a question, not its truth or the speaker's intent, and its output should not be read as a verdict on any politician. The LLM "Presidents" of Section 7 are role-play prompts, not models of real people's views. All data are the public interview transcripts released for the shared task.

Limitations, Challenges and Future Work

One small development set. Every system comparison is on 308 dev items, because the test labels were never released and the evaluation server has closed. Differences between systems therefore carry wide intervals, and parameters fitted on dev (by nested cross-validation) use dev labels. *Next step:* score the frozen test predictions in the repository as soon as the organisers release the test labels, and re-run the instrument, unchanged, on a second evasion corpus such as parliamentary question periods.

Validity of the instrument on LLM text. The instrument was trained on presidential speech, and our check on LLM answers is label-free, so it cannot rule out confident errors. *Challenge:* the registered human annotation of LLM answers was dropped for lack of time. *Next step:* have two annotators label a stratified sample at the clarity level, blind to model and condition, and use the instrument only if its agreement with them is within 0.10 of their agreement with each other, as registered.

Breadth of the LLM audit. The audit covers two model families, one of which shares a family with the 27B judge. The role effect was not a registered hypothesis. Labelling N and P with one instrument removes a question prior that acts alike on every answer, not one that interacts with answer style, and we had no working question-only control. The role odds ratio was not adjusted for answer length. *Challenge:* two further families (Llama-3.1, OLMo-2) failed to load on the shared GPU because of context-length and memory settings. *Next step:* re-run them with their own GPU and context limits, adjust the role odds ratio for answer length, and register the role effect as a hypothesis for that replication. Role play is also not politics: an LLM has no career at stake, so "role" here means the prompt, and conclusions about real politicians would need human studies.

Reproducibility of LLM decisions. The adjudicator's decisions depend on the serving stack, the GPU and the option order, not only on the weights and the item, and we never repeated a run, so its own run-to-run noise is unmeasured. *Next step:* repeat the 27B's dev run several times, pin the serving software and hardware in the frozen file, and report decision agreement across repeats alongside every result that uses an LLM.

Cost measurements. GPU time was measured with batched inference over the 308 dev items on one GPU type. Per-item cost depends on batch size and hardware, and our reduced-precision runs used unfused FP8 kernels, so they measure accuracy under quantization but not the speed a production kernel would give. The winning system's cost is known only in calls and tokens. *Next step:* serve M1 through fused FP8 kernels, retrain the adapter on the 4-bit base (QLoRA), as our registered trigger requires, score the 27B on every item NF4 defers so its cascade comparison is complete, and report latency at batch size one alongside throughput.

Information the system never sees: the question. Every module reads the answer, yet the role study suggests that the question carries evasion signal of its own: models answering as the President, without the real answer, answered directly the questions the Presidents answered directly (Section 7, exploratory). *Next step:* add a question-side prior, module M5 in Figure 2. For the items M2 defers, two open LLMs answer the question in the President role, the frozen instrument labels their answers, and the share that evade enters the fusion as an extra term. Its weight would be fitted on the training rows (whose P-role answers must first be generated) and its prediction registered before dev is scored. We expect a small effect on the items M1 is unsure about, mainly on the evade-or-not decision. Appendix B gives a preliminary check. *Challenge:* it costs two extra generations per deferred item and relies on the instrument being valid on LLM text.

Labels and classes. Training items carry one annotator's label, so the models cannot learn how disputed an item is, and the any-annotator metric credits labels that most annotators rejected. Partial/half-answer and General remain the weakest classes. *Next step:* train on soft labels where several annotations exist, and add targeted contrast examples for the weakest classes. We also cannot rule out that

the pretraining data of the LLMs include the public interview transcripts.

Process challenges. Shared hardware forced two deviations from our plans (a different GPU type for E25, and two LLM families dropped), both logged before results were seen. A fusion bug in our own analysis was found and corrected in the log. Every analysis now reproduces its saved predictions before it reports a number.

References

R. Michael Alvarez and Jacob Morrier. 2025. Measuring the quality of answers in political Q&As with large language models. *Political Analysis*.

Lingjiao Chen, Matei Zaharia, and James Zou. 2023. FrugalGPT: How to use large language models while reducing cost and improving performance. *arXiv preprint arXiv:2305.05176*.

Steven E. Clayman. 2001. Answers and evasions. *Language in Society*, 30(3):403–442.

Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. QLoRA: Efficient finetuning of quantized LLMs. In *Advances in Neural Information Processing Systems 36 (NeurIPS)*.

Pengcheng He, Jianfeng Gao, and Weizhu Chen. 2023. DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing. In *The Eleventh International Conference on Learning Representations (ICLR)*.

Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. 2015. Distilling the knowledge in a neural network. *Preprint*, arXiv:1503.02531.

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. LoRA: Low-rank adaptation of large language models. In *The Tenth International Conference on Learning Representations (ICLR)*.

Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, Lélio Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timothée Lacroix, and William El Sayed. 2023. Mistral 7B. *arXiv preprint arXiv:2310.06825*.

Aditya Krishna Menon, Sadeep Jayasumana, Ankit Singh Rawat, Himanshu Jain, Andreas Veit, and Sanjiv Kumar. 2021. Long-tail learning via logit adjustment. In *The Ninth International Conference on Learning Representations (ICLR)*.

Wisarut Peerachaidecho and Attapol Rutherford. 2026. ChulaNLP at SemEval-2026 task 6: A hybrid BERT-LLM framework for political response clarity and evasion detection. In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 1874–1881, San Diego, California, USA. Association for Computational Linguistics.

Lingling Shi, Haoyu Jin, Shiquan Wang, Fang Yu, Shuangyong Song, and Xuelong Li. 2026. TeleAI at SemEval-2026 task 6: A confidence-aware multi-stage reasoning framework with chain-of-thought. In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 1591–1609, San Diego, California, USA. Association for Computational Linguistics.

Konstantinos Thomas, Giorgos Filandrianos, Maria Lymperaiou, Chrysoula Zerva, and Giorgos Stamou. 2024. "I never said that": A dataset, taxonomy and baselines on response clarity classification. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 5204–5233, Miami, Florida, USA. Association for Computational Linguistics.

Konstantinos Thomas, Giorgos Filandrianos, Maria Lymperaiou, Chrysoula Zerva, and Giorgos Stamou. 2026. SemEval-2026 task 6: CLARITY – unmasking political question evasions. In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 3704–3715, San Diego, California, USA. Association for Computational Linguistics.

An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, et al. 2025. Qwen3 technical report. *Preprint*, arXiv:2505.09388.

Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. 2024. Large language models are not robust multiple choice selectors. In *The Twelfth International Conference on Learning Representations (ICLR)*.

A Experiment Ledger

Table 4 lists every experiment in the order it ran. The hypotheses, decision rules and amendments are in the repository's pre-registered log.

B Implementation and the Frozen Instrument

M1 and the students. Qwen3-8B-Base in bf16 with LoRA ($r = 16$, $\alpha = 32$, dropout 0.05) on all attention and feed-forward projections and a new nine-way head on the last non-padding token. Training uses AdamW at 10^{-4}, a cosine schedule with 5% warm-up, effective batch 16 and 12 epochs, keeping the epoch with the best slice macro-F1. The input is "Sub-question / Full question" ($\leq$256 tokens) and the answer, $\leq$1,024 tokens in all, with a fixed closing question. Students use the same recipe with the soft-target cross-entropy. Cross-fitted candidates for distillation come from five models trained on four fifths of train each.

M3. Qwen3.8-27B-FP8 served with vLLM, thinking disabled, greedy decoding, one output token, top-20 log-probabilities. The prompt holds the task, definitions of all nine labels, a confusion guide for the most-confused pairs (Explicit–Implicit, Implicit–General, Explicit–General, Deflection vs Implicit/General, Dodging–Deflection), the item in full and the four lettered candidates. q is averaged over four cyclic option orders.

The frozen instrument. All settings are in `code/experiments/e25/frozen_system.json`: M1's ten seeds, slice calibration $\tau_C = 0.20$ and temperature $T = 2.15$, top-4 candidates, and deferral iff $u \geq 0.638$ (the slice's 70% quantile). Kept items use $\tau_1 = 1.15$ and deferred items $\lambda = 0.5$, $\tau = -0.55$, fitted once on dev with the averaged records. The same procedure under nested CV gives 0.650 at split 0 and 0.639 over ten splits. Run on dev with the E23 records, it reproduces the frozen predictions exactly.

E25b prompts. All conditions are one user message asking for spoken prose, no lists, at most 250 words, with greedy decoding and at most 400 new tokens. N casts the model as a neutral analyst who is not the President, and P as the President in that interview, on its date. F appends to P's conversation "With respect, you haven't answered the question." and the sub-question. B adds the President's real answer to P as private notes. N, P and B are generated once per interview turn and shared by its sub-questions, as the President's one answer is.

Cost measurement. GPU-s per item is the wall time of one batched pass over the 308 dev items divided by 308, with the GPU used by that job alone. To keep every comparison like for like, all reported times come from one session on one GPU (an RTX PRO 6000): an M1 seed (bf16, batch 16) takes 0.040 GPU-s per item and a 27B call 0.098–0.101 per order. A system's time is K seeds plus its realised 27B calls. Ratios computed from an earlier measurement on another card of the same model differ by at most 0.02$\times$ (`docs/raw/cost_same_gpu.txt`).

Preliminary check for the question-side prior. Written after the E25 results, on the 308 dev items only (`docs/raw/question_prior_check.txt`). Of the two P-role LLMs, the share whose answer the instrument labels evasive was 0, 0.5 or 1 for 101, 95 and 112 questions, and the President evaded on 61%, 72% and 88% of them. As a score for "the President evaded" it reaches an AUROC of 0.661 (M1: 0.880, or 0.872 cross-validated) and correlates 0.31 with M1. In a cross-validated logistic regression it lifts M1's AUROC from 0.872 to 0.875 on all items and from 0.722 to 0.746 on the 111 items M1 is least sure of. This motivates a registered test and is not evidence for the module.

Code. Code, the pre-registered log and the raw output behind every number: https://github.com/IamSid44/CLARITY-Unmasking-Political-Question-Evasions.

Exp.	Question	Key result (dev S2)	Verdict	Led to
E0–E3	Where does DeBERTa fail?	gold in top 3: 90%	–	decision rules
E4	Decision rule?	LA 0.365 $\rightarrow$ 0.438	held	LA in every system
E5, E6	Taxonomy routing, definition re-ranker?	0.369, 0.354 vs 0.365	failed	knowledge, not structure
E8, E9	Rebalanced loss, hierarchical head?	0.379, 0.370 vs 0.438	failed	train longer
E10, E11	16 epochs, full question?	+0.047, +0.022 per model	held	ten paired seeds
E12	Gate, specialists, pair experts, soups?	$-0.059 \ldots -0.117$	failed	change the backbone
E13	Is the gap knowledge?	+0.092 per model, 10/10	held	LLM classifier
E14	M1 at 10 seeds, can K shrink?	0.598, K=5: 0.555	held / failed	M3, fused
E15	Does a 27B chooser help?	0.669 (after bug fix)	held	accuracy point
E16	Prompt parts, cost, smaller ensembles?	seeds 92% of time, smaller ensembles gain less	–	distil, test the headline
E17, E18	One distilled student?	0.608 with M3 at 0.25$\times$ M1's time	mixed	distilled point
E19	Controls on the cascade?	train rows +0.014 [+0.002, +0.026]	held	routing and fusion carry it
E20	Learned router?	loses 4/4 budgets	failed	confidence routing
E21	Test transcripts differ?	15.6% labels change	held	normalised test
E22	Cache, ladder, relevance?	cache 0.49$\times$ but 9.4% flips, 0.600 $<$ 0.625, $\rho = +0.095$	failed / held	E24
E23	Final re-scoring, one setup	0.650 / 0.608	–	frozen instrument
E24	What is the label a property of?	siblings 72% vs 43%	–	evasion is relational
E25a	Question read causally?	CR 0.253 $\rightarrow$ 0.013	held	audit is valid
E25b	Question or speaker?	OR 1.18 (N), 2.74 (P)	failed / expl.	question $\times$ role
E25c, d	Early or frozen shortcut?	66% at half, probe 0.397	failed	no shortcut
E26	Cheaper by lower precision?	FP8 0.83$\times$, M1 -0.045, instrument -0.003	failed / held	fused kernels, QLoRA
Pre-checks rejected in E24 on existing outputs				
P3	Route on seed disagreement?	49 vs 53 of 115 errors caught	failed	–
P4	Decode siblings jointly?	$\pm$0.01, sign flips	failed	–
P5	Untuned metric-optimal rule?	0.594 vs 0.598	failed	–
P6	One debiased option order?	agreement 0.85 $\rightarrow$ 0.79	failed	–
P7	Conformal sets as router?	0.564–0.617 vs 0.641–0.663	failed	–

Table 4: Every experiment, in order. E0–E13 precede the mid-evaluation and E14–E26 follow it. LA is logit adjustment, CR the Clear-Reply share, OR the odds ratio and expl. exploratory.

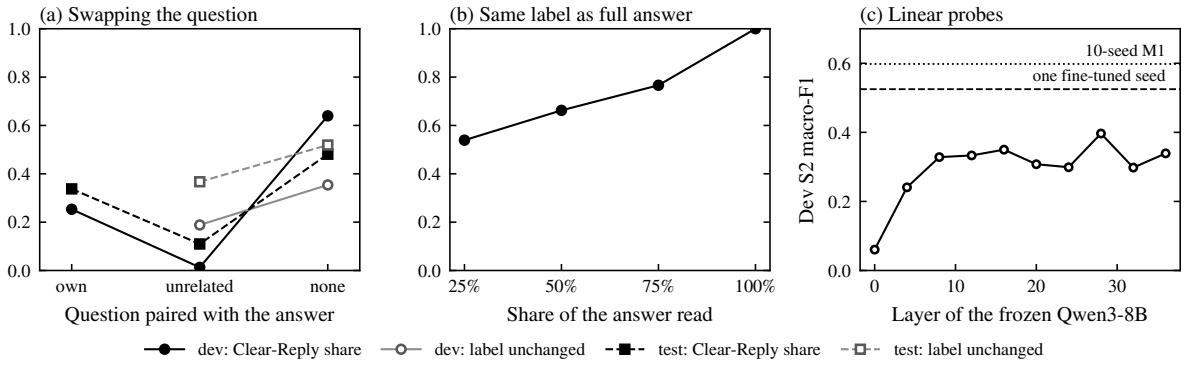


Figure 9: E25a, c and d. (a) Each answer paired with its own question, a question from another interview, or none. "Label unchanged" is the share of items keeping their own-question label. (b) Agreement of the label on the first 25/50/75% of the answer with the full-answer label (dev). (c) Linear probes on every fourth layer of the frozen Qwen3-8B-Base.