

Knowledge over Structure: A Controlled Study of Response Clarity Classification in Political Interviews

Vidvathama R Siddarth Gottumukkula Sanjana Reddy Vonteri Shashikanta Sahoo
2024122002 2023102040 2026901007 2026900007

Team Nier_ANLP
International Institute of Information Technology, Hyderabad

Abstract

SemEval-2026 Task 6 (CLARITY) asks how a politician’s answer responds to an interview question. Subtask 2 has nine evasion types, and each type fixes one of the three clarity levels of Subtask 1. We set out to find what limits single-pass classifiers on this task, changing one thing at a time and writing down our predictions before every run. A DeBERTa-v3-large classifier built this way reaches 0.384 multi-reference macro-F1 on Subtask 2 per model, and 0.405 as a ten-seed system with a one-parameter logit adjustment. Re-ranking, three kinds of hierarchy, rebalanced losses and model soups all failed to improve on it, and the error analysis points at the model rather than at label noise. When we swap only the backbone for Qwen3-8B fine-tuned with LoRA, keeping the data, input, loss and model selection fixed, Subtask 2 rises by 0.092 and Subtask 1 by 0.095 on every one of ten paired seeds, and the ten-seed system reaches 0.543 and 0.746 (+0.136 on Subtask 2, 95% CI [0.054, 0.224]). The gain is spread over most classes instead of sitting on the label boundaries we had expected. All numbers are on the development set, because the test labels were never released.

1 Introduction

Politicians often answer the question they would have liked to be asked. [Thomas et al. \(2024\)](#) turned this observation into a classification task with a two-level taxonomy of response clarity and evasion, and SemEval-2026 Task 6 ([Thomas et al., 2026](#)) made it a shared task. Subtask 2 asks for one of nine evasion types, and Subtask 1 for one of three clarity levels, which follow directly from the nine types.

The competition was won by pipelines that call a large language model (LLM) several times per item ([Shi et al., 2026](#); [Peerachaidecho and Rutherford, 2026](#)). Fine-tuned encoders, by contrast, stopped at about 0.50 macro-F1 on Subtask 2 no matter

how large they were or how many were ensembled ([Thomas et al., 2026](#)). A leaderboard cannot tell us *why* the encoders stop there, and the possible reasons predict different things. If the labels are too noisy to learn, models should fail where annotators disagree. If the decision rule is wrong for the metric, a better rule should help. If the model lacks the knowledge needed to read an answer against its question, a larger pretrained model should help under the same protocol. And if the meaning of the labels cannot be learnt from 3,448 examples, giving the model definitions should help.

We test these explanations with a deliberately controlled single-pass classifier (Figure 1). Every change is made on its own against a named control, every claim rests on ten paired seeds, our predictions were written down before each run, and nothing is chosen on the development set. This gives four contributions.

- A DeBERTa-v3-large system (§4) whose two real improvements, longer training and the full journalist question, are separated from a series of documented failures: re-ranking, hierarchical models, rebalanced losses, per-class thresholds and model soups.
- Evidence that the remaining gap lies in the model and not in label noise or the decision rule (§6).
- A one-change test of the knowledge explanation (§5). Replacing the backbone with Qwen3-8B-Base fine-tuned with LoRA ([Yang et al., 2025](#); [Hu et al., 2022](#)) improves both subtasks on all ten seeds and lifts the ten-seed system from 0.405 to 0.543 on Subtask 2 and from 0.648 to 0.746 on Subtask 1.
- An analysis of where that gain falls, which shows that our own registered prediction was wrong (§6).




Figure 1: The classification system. One backbone (DeBERTa-v3-large, fully fine-tuned, or Qwen3-8B-Base with LoRA) reads the sub-question, the journalist’s full question and the answer, and a new linear head gives a distribution over the nine evasion types. A system averages ten seeds and applies logit adjustment, which divides by the train class prior π_c raised to τ . The Subtask 1 clarity level follows from the Subtask 2 label through the fixed taxonomy.

Because the test labels were never released, all comparisons are on the 308-item development set, and we report seed spreads, paired differences and bootstrap intervals throughout.

2 Problem Statement, Data and Evaluation

Problem. SemEval-2026 Task 6 (Thomas et al., 2026) uses the question–answer pairs of Thomas et al. (2024), taken from U.S. presidential interviews. Each item is a triple $x = (q, Q, a)$. Here q is a sub-question (median 15 tokens), Q is the journalist’s full turn from which it was split (median 62 tokens), and a is the president’s entire answer (median 266 tokens). Let $\mathcal{Y}_9$ be the nine evasion types, $\mathcal{Y}_3$ the three clarity levels, and $g : \mathcal{Y}_9 \rightarrow \mathcal{Y}_3$ the fixed taxonomy map in Table 1. Subtask 2 (S2) asks for $f(x) \in \mathcal{Y}_9$, the way that particular sub-question was answered, and Subtask 1 (S1) asks for a level in $\mathcal{Y}_3$. Like the top-ranked systems (Shi et al., 2026; Peerachaidecho and Rutherford, 2026), we learn f and predict $g(f(x))$ for S1. We call a classifier *single-pass* if it runs one model once per item, and the question of this paper is what limits such a classifier on this task.

Data. The training set has 3,448 items with one label each. The development (dev) set has 308 items labelled by three annotators, of which 125 are unanimous, 150 split two to one and 33 get three different labels, so the reference set G_i of an item holds one to three types. The 237 test items have two annotators each, but their labels were never released and the evaluation server has closed, so every comparison in this paper is on dev. Two properties of the data shaped our modelling. First, 69% of training rows share their answer with another sub-question, and 71.7% of those shared answers carry different labels, so the model has to

Evasion type	Clarity level	Train %	Dev sets
Explicit	Clear Reply	30.5	115
Implicit	Ambivalent	14.2	99
Dodging	Ambivalent	20.5	96
General	Ambivalent	11.2	113
Deflection	Ambivalent	11.0	51
Partial/half-answer	Ambivalent	2.3	14
Declining to answer	Clear Non-Reply	4.2	17
Claims ignorance	Clear Non-Reply	3.5	15
Clarification	Clear Non-Reply	2.7	4

Table 1: The taxonomy. Each evasion type determines a clarity level. Train % is the share of the 3,448 training labels, and Dev sets is the number of the 308 dev items whose reference set contains the type.

work out which part of a long answer responds to which sub-question. Second, the answers are long, and for 101 of the 308 dev items the sub-question and answer together exceed 512 DeBERTa tokens.

Metric. Both subtasks are scored by macro-F1 over all declared classes, so a class that is never predicted contributes $F_1 = 0$. S1 uses plain macro-F1. S2 uses *multi-reference macro-F1*, in which a prediction $\hat{y}_i$ is a true positive for class $\hat{y}_i$ if $\hat{y}_i \in G_i$ and a false positive otherwise, while a miss charges one false negative to every member of G_i . If every prediction is acceptable, every class predicted at least once has $F_1 = 1$, and

$$\text{macro-F1} = \frac{|\{\hat{y}_1, \dots, \hat{y}_N\}|}{9}. \quad (1)$$

The score therefore rewards two separate things, landing in the reference set (the *in-set rate*) and *coverage*, naming each class at least once, so that Clarification (4 dev reference sets) counts as much as Explicit (115). The floors show the trade-off. Always predicting Explicit lands in-set for 37.3% of dev items but scores only 0.060 on S2, uniform

guessing scores 0.129, and TF-IDF with logistic regression scores 0.257 (on S1, 0.136, 0.299 and 0.445).

Evaluation protocol. We choose the training epoch on the *train slice*, a fixed stratified 10% of train (345 rows) held out from training, and never use dev to choose anything. Every system claim rests on *ten paired seeds*, with the same seed numbers in both arms, reported as mean and standard deviation together with the paired difference, its *t* statistic and the number of seeds that improve. Each variant is first screened on three seeds against a bar fixed in advance, a gain of at least +0.015 with at least two of the three seeds up. A *single model* is one seed scored by argmax with no tuning on dev, and a *system* averages the seeds’ probabilities and applies logit adjustment (Menon et al., 2021),

$$\hat{y}(x) = \arg \max_c \frac{\bar{p}(c | x)}{\pi_c^\tau}, \quad (2)$$

where $\bar{p}$ is the mean of the seeds’ probabilities, π_c is the frequency of class c in train, and τ is the only fitted parameter. Since τ is fitted on dev, we choose it by 5-fold nested cross-validation (grid from -1 to 2 , step 0.05) and score the pooled held-out predictions, reporting the first split and, where given, the mean over ten splits. Systems are compared with a bootstrap over dev items with the seeds held fixed (2,000 resamples), and because the test set has two annotators we also rescore dev against each of its three two-annotator reference sets. Before each experiment, we wrote its hypotheses and predicted outcomes into an experiment log.

3 Related Work

Task and top systems. Thomas et al. (2024) introduced the clarity and evasion taxonomy over U.S. presidential interviews. They found that predicting the nine-way label and mapping it up the hierarchy worked better than predicting clarity directly. In the shared task (Thomas et al., 2026), the organisers’ fine-tuned Llama-70B scored 0.82 (S1) and 0.57 (S2) on test. The organisers report that LLM prompting and hierarchical use of the taxonomy were the most effective strategies, and that fine-tuned encoders levelled off near 0.50 on test S2. TeleAI (Shi et al., 2026) ranked first (test 0.89 S1, 0.68 S2) with three confidence-routed chain-of-thought stages on DeepSeek-V3, whose prompts include boundary examples mined from the training confusion matrix. Their pipeline makes 1.94

calls per item on average and reaches 0.617 (S2) and 0.812 (S1) on dev. In the same paper, a directly fine-tuned Qwen2.5-7B reached 0.495 S2 on dev. ChulaNLP (Peerachaidecho and Rutherford, 2026), second on S2, passed the top five labels of a fine-tuned DeBERTa-large to few-shot Kimi-K2 and scored 0.52 S2 on dev, against 0.46 for the encoder alone. Since the test labels were never released, we compare with these systems on dev.

Other participating systems. PFW (Tamsal and Rusert, 2026) averaged the logits of 50 DeBERTa models per architecture (5 folds $\times$ 10 seeds) and reached 0.76 (S1) and 0.50 (S2) on test. Three post-hoc techniques that raised their out-of-fold scores lowered their test scores by 0.02–0.10, which mirrors how per-class decision weights overfit dev for us. SilkPeak (Tetakali and Tetakali, 2026) trained a single DeBERTa-v3-large with joint clarity and evasion heads (0.76 S1 on test) and found it ahead of hierarchical bi-encoders and a LoRA-tuned LLaMA-3-8B. CLaC (Turk et al., 2026) found prompted LLMs ahead of fine-tuned encoders and placed third on S2 (0.59). ttda704 (Tran Tan and Dinh, 2026) found structured chain-of-thought prompting (0.515 S2) ahead of parameter-efficient fine-tuning of Qwen3 models, and Team Evaluators (Nuthakki et al., 2026) reached 0.54 S2 with LoRA-tuned open LLMs. These systems change several components at once, whereas we change one at a time.

Methods we build on. Logit adjustment (Menon et al., 2021) subtracts scaled log class priors from the logits after training. Balanced Softmax (Ren et al., 2020) makes the same prior correction inside the loss, and focal loss (Lin et al., 2017) down-weights easy examples. Our analysis of multiple annotators treats their disagreement as signal (Uma et al., 2021). Our encoder is DeBERTa-v3-large (He et al., 2023), and our LLM classifier is Qwen3-8B-Base (Yang et al., 2025) fine-tuned with LoRA (Hu et al., 2022) and used as a single-pass classifier.

4 The Encoder Track

Model and training. We fine-tuned all 435M parameters of DeBERTa-v3-large (He et al., 2023) as a 9-way classifier, with the [CLS] state passed through a dense tanh pooler and a new linear layer over the nine labels. The baseline input pairs the sub-question (at most 96 tokens) with the answer, cut from the right at 512 tokens. The final input

Configuration	Single model		System	
	S2	S1	S2	S1
Baseline, 8 ep.	0.315 $\pm$ 0.040	0.576 $\pm$ 0.030	0.428	0.601
+ 16 epochs	0.362 $\pm$ 0.026	0.601 $\pm$ 0.016	0.400	0.620
+ full question	0.384 $\pm$ 0.030	0.614 $\pm$ 0.027	0.405	0.648
All of train [†]	0.416 $\pm$ 0.020	0.614 $\pm$ 0.016	0.489	–

Table 2: DeBERTa-v3-large on dev. Single model is the mean $\pm$ sd over seeds 0–9. System is the 10-seed mean probability with logit adjustment (τ by nested CV). Rows 2 and 3 each add one change, and row 3 is the final system. [†]All of train, last epoch, seeds 0–4 only, so its system is a 5-seed number and not comparable.

puts “Sub-question: ... Full question: ...” first, with up to 256 tokens, and raises the context to 1,024 tokens. We added the full question because 69% of training rows share their answer with another sub-question, and a model that cannot see the other sub-questions has no way to tell which part of the answer responds to which. Training used plain cross-entropy, AdamW at 10^{-5} with layer-wise decay 0.95, a cosine schedule and batch size 16 (Appendix A), and the epoch was chosen on the train slice. A seed fixes the head initialisation and the data order, so seeds 0–9 are paired across configurations.

What improved the model. Two changes improved single models (Table 2, Figure 2). Training for 16 epochs instead of 8 raised S2 by +0.047 ($t = 3.4$, 9 of 10 seeds up) and S1 by +0.026 ($t = 2.7$, 8 of 10). Adding the full question gave a further +0.022 on S2 ($t = 2.0$, 7 of 10) and +0.012 on S1 ($t = 1.3$, 6 of 10). Together the two changes add +0.069 on S2 and +0.038 on S1, each on 9 of 10 seeds. Most of the gain came from longer training, which also helped the classes the baseline handled worst. General, for example, rose from 0.119 to 0.295 per model (seeds 0–4).

At first, undertraining looked like a negative result. With 8 epochs the full question scored below the baseline (0.282 ± 0.043 against 0.337 ± 0.044 , seeds 0–4), but its final training loss (1.32–1.50, against 0.57–1.09) showed it had not finished fitting. Its seeds were less confident (mean top probability 0.445 against 0.594) and leaned on the class prior, and at 16 epochs the same input gave the best single model.

From model to system. Better single models did not give a better S2 system. Logit adjustment lifted the baseline ensemble by +0.109 (to 0.428) but did




Figure 2: Single DeBERTa-v3-large models on dev for the three configurations: (a) S2 multi-reference macro-F1, (b) S1 macro-F1. Open circles joined by thin lines are individual seeds, paired across configurations, and filled squares are the mean over the 10 seeds, with error bars of ± 1 standard deviation.

nothing for the final ensemble (0.409 to 0.405). The S2 difference of -0.022 between the two systems has a 95% bootstrap interval of $[-0.101, +0.061]$, while on S1 the final system is ahead (0.648 against 0.601). We think, although we have not tested it, that longer training and logit adjustment fix the same fault, the pull towards frequent classes. We also learnt that five-seed systems are unreliable on 308 items. The baseline system scored 0.438 on seeds 0–4 and 0.356 on seeds 5–9, a spread larger than most of the effects we wanted to measure, so all our system claims use ten seeds. Training on all of train for 16 epochs and keeping the last epoch gained +0.029 on S2 per model on seeds 0–4 (4 of 5 up) and lost 0.017 on S1, which with five seeds is not enough to rank it.

What did not help. Seven further ideas, each tested against a named control, left the system where it was. A second encoder that re-ranked the baseline’s top three labels using their definitions scored 0.354 against 0.365 for the ensemble it re-ranked, which suggests that learning what each label means from 3,448 examples is harder than classifying them. Balanced Softmax with focal loss (Ren et al., 2020; Lin et al., 2017) over-corrected towards rare classes (Explicit F1 fell from 0.679 to 0.319) and used up the gain that logit adjustment would otherwise have given. Three hierarchy designs (routing through the clarity level, a factorised head, and a Non-Reply gate with branch specialists) did not beat the flat softmax, which already scores all nine labels jointly. Boundary experts for the three most-confused pairs changed S2 by only

+0.002, because they rarely overturn a decision the flat model makes from the same input. A uniform soup of the ten models scored 0.267, since seeds with different head initialisation and data order do not average well. Nine per-class decision weights overfit the 308 items, and an ensemble mixing both inputs scored 0.403 against 0.405. Since nothing trained on the same 3,448 examples moved the system, §5 tests the backbone itself.

5 Is the Gap Knowledge? An 8B LLM Classifier

The encoder track and the error analysis (§6) show that the model, and not label noise, is what fails. That leaves two explanations. The encoder may lack *knowledge* that a 0.4B-parameter model never acquired, or it may lack *label meaning*, which a classifier has to infer from about 3,400 examples. A larger pretrained classifier tests the first explanation with nothing else changed. Before any run we registered three predictions. The gain would pass the screening bar, the single-model S2 mean would lie in 0.43–0.50 with S1 rising by +0.02 to +0.05, and the gain would fall mostly on the commitment boundary (Explicit $\leftrightarrow$ Implicit, General $\leftrightarrow$ Implicit/Explicit) and more on agreed than on contested items.

One change: the backbone. We took the best encoder configuration (sub-question, full question and answer, 16 epochs) and replaced DeBERTa-v3-large with Qwen3-8B-Base (Yang et al., 2025). The model is adapted with LoRA (Hu et al., 2022) of rank 16 ($\alpha = 32$, dropout 0.05) on all linear projections of the attention and feed-forward blocks, and a new 9-way linear head, trained in full, reads the hidden state of the last token (Figure 1). The input text and token budgets are the encoder’s, with one fixed closing question appended (“How does the answer respond to the sub-question?”) so that the token the head reads comes after the whole question and answer in every item. Every run first checks that an item’s logits are the same alone and inside a right-padded batch, which guards against a head that reads a padding token. Training used AdamW at 10^{-4} for 3 epochs with an effective batch of 16, and the learning rate and epoch count change only because the backbone does. Everything else stays identical, namely the same 3,103 training rows, the same 345-row train slice (asserted identical in code), plain cross-entropy and the epoch with the best slice macro-F1. Seeds 0–

9 are paired with the DeBERTa runs of the same seeds, and Appendix A lists all hyperparameters and the compute.

Results at ten paired seeds. The 3-seed screen passed the bar (+0.081 on S2, 3 of 3 seeds up), so we extended the run to ten seeds (Table 3). Per model, Qwen improves S2 by +0.092 ($t = 7.12$) and S1 by +0.095 ($t = 8.48$) on all ten seeds, which is more than the whole encoder track gained (0.315 to 0.384). The system gain holds on the 2-annotator reference sets (0.524 against 0.382). All the registered predictions about size held, with the S1 gain even exceeding its range, but the prediction about *where* the gain falls turned out wrong (§6). Among published dev results (Figure 3), the system lies above TeleAI’s fine-tuned Qwen2.5-7B (0.495) and ChulaNLP’s pipeline (0.52) and below TeleAI’s multi-call pipeline (0.617), though these come from other protocols.

Two single-change follow-ups. In the 3-epoch runs the slice F1 rose at every epoch, and the last epoch was always the one selected. Under a cosine schedule this proves little, because the schedule itself favours the last epoch, so only a longer schedule can test for undertraining. With 12 epochs (seeds 0–2) the slice selected epochs 9, 11 and 8, and the slice F1 rose by +0.084 on all three seeds, so the rule we had registered adopted the longer schedule. Training loss fell to about 0.02 without the slice F1 declining, and dev S2 was 0.543 ± 0.015 , which we report but did not use for the choice. Training on all 3,448 rows gave +0.017 per model (6 of 10 seeds up), but the system difference has a 95% interval of $[-0.026, +0.092]$, and S1 fell from 0.746 to 0.711. As with the encoder, the 3-seed screen (+0.039) had overstated the gain.

6 Analysis

We analyse the final DeBERTa 10-seed ensemble (argmax) and compare it with the paired Qwen runs of §5, all on dev.

The model, not label noise, fails. Of the ensemble’s 308 predictions, 137 (44.5%) fell outside the item’s reference set, and 53 of these were on items where all three annotators agreed. Annotator disagreement is real and informative (Uma et al., 2021), but it does not explain the gap. Scored the way the test set is scored, each annotator against the other two reaches 0.643–0.766 on S2 (mean 0.684), while the ensemble reaches 0.380–0.381.

Dev macro-F1	Qwen3-8B + LoRA	DeBERTa-v3-large	Δ	Seeds up	95% CI
Single model, S2	0.476 $\pm$ 0.043	0.384 $\pm$ 0.030	+0.092	10/10	
Single model, S1	0.709 $\pm$ 0.031	0.614 $\pm$ 0.027	+0.095	10/10	
Single model + LA, S2	0.538 $\pm$ 0.043	0.389 $\pm$ 0.019	+0.149	10/10	
System, S2	0.543 (0.549)	0.405 (0.412)	+0.136		[+0.054, +0.224]
System, S1	0.746	0.648	+0.098		
System, S2, 2-annotator sets	0.524	0.382	+0.142		
Follow-up (one change)	Variant	Qwen, 3 epochs	Δ	Seeds up	95% CI
12 epochs, single model, S2 [‡]	0.543 $\pm$ 0.015	0.462 $\pm$ 0.080	+0.082 [†]		
All of train, single model, S2	0.492 $\pm$ 0.045	0.476 $\pm$ 0.043	+0.017	6/10	
All of train, system, S2	0.575 (0.562)	0.543 (0.549)	+0.028		[−0.026, +0.092]

Table 3: The backbone swap (top) and two follow-ups (bottom) on dev. Single model is the mean $\pm$ s.d. over seeds, and Δ the paired mean difference. LA is logit adjustment. System is the 10-seed ensemble with LA, τ chosen by nested CV (in parentheses, the mean over 10 CV splits). CIs are bootstraps over dev items, and the 2-annotator sets follow the test regime. [‡]Seeds 0–2. [†]Reported only, since the 12-epoch schedule was adopted on the train slice.




Figure 3: Dev S2 multi-reference macro-F1 of the DeBERTa-v3-large baseline, DeBERTa with the full question and 16 epochs, Qwen3-8B-Base with LoRA on the same input (3 epochs), and Qwen trained on all of train. (a) Single models on 10 paired seeds: open circles joined by thin lines are individual seeds, and filled squares are means with ± 1 standard deviation. (b) The corresponding 10-seed systems (ensemble + logit adjustment, τ by nested CV, bar height the mean over 10 CV splits), with published dev results as horizontal lines: TeleAI’s multi-call pipeline (0.617) and fine-tuned Qwen2.5-7B (0.495) (Shi et al., 2026), ChulaNLP’s DeBERTa top-5 + Kimi-K2 pipeline (0.52) (Peerachaidecho and Rutherford, 2026), and a human annotator scored against the other two (0.684).

Commitment, not only topic. The largest error pairs (reference $\rightarrow$ prediction) are Implicit $\rightarrow$ Explicit (25), General $\rightarrow$ Implicit (20) and General $\rightarrow$ Explicit (20). All three concern how committed an answer is, more than whether it stays on topic. Length matters too. The 20 inputs that are still cut at 1,024 tokens land in-set only 20% of the time (0.580 for the rest), and the shorter half of dev scores 0.400 on S2 against 0.283 for the longer half.

A typical error shows how much an answer can depend on world knowledge (Figure 4). Asked whether he would campaign against Senator Lieberman, the speaker replied “*I’m going to stay out of Connecticut.*” Seeing this as a near-answer requires knowing that Lieberman is a senator from Connecti-

cut, which is our reading of the item rather than something we measured, and errors of this kind are what led us to the knowledge hypothesis. Qwen ranks Dodging first, but logit adjustment then picks the rarer Declining to answer. A prior correction that raises macro-F1 over the whole set can still cost an individual item.

Where the 8B classifier gains. Table 4 shows that our registered prediction was wrong. Per-class F1 rises in 7 of the 9 classes, most in Claims ignorance (+0.403), Declining (+0.220), Dodging (+0.103) and Deflection (+0.089). The commitment classes move least (General +0.059, Explicit +0.057, Implicit +0.040), and neither model ever predicts Partial/half-answer. The in-set rate rises

Sub-question: Would you campaign against Senator Joe Lieberman on Iraq?

Full question: ... would you campaign against Senator Joe Lieberman, whose Republican candidate may support you, but he supports you too, on Iraq?

Answer: I'm going to stay out of Connecticut.

Annotators: Dodging, Dodging, Implicit (S1: Ambivalent)

	DeBERTa system	Qwen system
$\bar{p}$, top 3	General .373 Implicit .282 Deflection .205	Dodging .327 Implicit .189 Declining .170
Ensemble argmax	General $\times$	Dodging $\checkmark$
τ (item's fold)	0.70	0.85
S2 after LA	General $\times$	Declining $\times$
S1	Ambivalent $\checkmark$	Clear Non-Reply $\times$

Figure 4: One dev item through both 10-seed systems ($\checkmark$: in the reference set, or for S1 the gold level). The item was fixed in advance as the error discussed in §6. Qwen ranks an acceptable label first, but logit adjustment divides by the train prior (0.205 for Dodging, 0.042 for Declining to answer) and moves the decision to the rarer class.

on unanimous items (+0.076) about as much as on contested ones (+0.105 and +0.048). The biggest gains are on the Non-Reply classes, which depend on recognising what kind of statement an answer is, while General stays the weakest frequent class (F1 0.322).

Ranking and the decision rule. Both models rank better than they decide. For Qwen (seeds 0–2), an acceptable label is the top prediction for 60.4% of items and in the top three for 93.1%, and the in-set rate rises from 0.40 to 0.87 across fifths of dev from least to most confident, so most remaining errors are a choice among a few candidates the model already ranks highly. Raw Qwen leans towards frequent classes. On seeds 0–2 it predicted General for only 24 items, although General is in 113 reference sets. Logit adjustment corrects this, adding +0.062 per Qwen model (0.476 to 0.538) against +0.005 for DeBERTa. Mixing Qwen and DeBERTa probabilities, with the weight fitted inside the nested CV, adds nothing (0.543 against 0.549 for Qwen alone).

7 Conclusion and Plan to the Final Evaluation

Our controlled encoder stops at 0.384 Subtask 2 macro-F1 per model (0.405 as a system), and the evidence says the limit is the model, not the labels or the decision rule. Swapping only the backbone

	DeBERTa	Qwen	Δ
<i>per-class F1</i>			
Explicit	0.672	0.729	+0.057
Implicit	0.403	0.443	+0.040
Dodging	0.484	0.587	+0.103
General	0.263	0.322	+0.059
Deflection	0.287	0.377	+0.089
Partial/half-answer	0.000	0.000	+0.000
Declining to answer	0.381	0.601	+0.220
Claims ignorance	0.296	0.700	+0.403
Clarification	0.667	0.524	−0.144
<i>in-set rate by distinct annotator labels</i>			
one ($n = 125$)	0.529	0.605	+0.076
two ($n = 150$)	0.493	0.598	+0.105
three ($n = 33$)	0.688	0.736	+0.048

Table 4: Where the backbone gain falls, for single models averaged over the same 10 seeds (dev S2). The registered prediction, a gain concentrated on the commitment boundary and on agreed items, does not hold.

for a LoRA-tuned 8B decoder lifts the system to 0.543 and 0.746 on all ten paired seeds. The gain is broad, which suggests that a larger pretrained model is simply better at reading what kind of response an answer is.

Uses beyond this task. For political interviews, the system flags how directly an answer addresses its question, which could help journalists and fact-checkers find non-answers in large archives. The same design should suit three settings we have not tested. In *earnings calls*, an executive's evasive reply to an analyst can carry information investors care about. In *parliamentary hearings*, oversight depends on whether officials actually answer the questions put to them. In *chatbot evaluation*, the key check is whether a reply resolves the question that was asked.

Goal. Our single-pass system is 0.074 below the winning pipeline on Subtask 2 (0.543 against 0.617 on dev) but makes no LLM call at all, whereas that pipeline makes 1.94 calls with about 7,300 input and 380 output tokens per item (Shi et al., 2026). Yet an acceptable label is in Qwen's top three for 93.1% of items (top five, 99.1%), and its confidence separates easy items from hard ones (in-set rate 0.87 in the most confident fifth, 0.40 in the least, §6), so an LLM should be needed on only some items. We therefore plan a confidence-routed cascade of four modules, described below, drawn in Figure 5 and scheduled in Table 5, each added as one change under the protocol of §2. We target at least 0.60 dev Subtask 2 with at most 0.3 LLM




Figure 5: The planned system for the final evaluation (dashed boxes are new modules, none run yet, see §7). LLM calls per item equal δ , against 1.94 for the winning system (Shi et al., 2026).

Dates	Module	Planned work	Exit criterion
Oct 5–11	M1	Twelve-epoch Qwen, seeds 3–9, cross-fitting, $C(x)$ and $u(x)$	Gold in $C(x)$ for $\geq 95\%$ of slice items
Oct 12–18	M3	Local LLM, prompt, exemplar index and candidate scoring	Beats M1 on deferred slice items
Oct 19–25	M2	Router and fusion, sweep of δ , ablation of each prompt part	Accuracy-vs-calls curve
Oct 26–29	M4	Distil M3 into M1, 4-bit inference, measured cost per item	Gain over control above seed spread
Oct 30–31	–	Freeze, final ten-seed and two-annotator scores, release	Final submission

Table 5: Timeline to the final evaluation (all planned). Each step adds one module to the system before it.

calls per item, about a sixth of the winner’s.

M1: calibrated candidate generator. M1 is the Qwen3-8B-Base classifier with LoRA, trained for the twelve epochs adopted on the train slice (§5). Its K seeds are averaged and logit-adjusted (Eq. 2) to give $\bar{p}(c | x)$, and since each seed costs one forward pass, K is our first cost lever. Instead of a single label, M1 returns a candidate set $C(x)$, the smallest set of top-ranked labels whose adjusted probability reaches a threshold. We calibrate that threshold on cross-fitted training probabilities (5 folds, so no item is scored by a model trained on it) so that $C(x)$ contains the gold label for 95% of held-out items. M1 also returns an uncertainty score $u(x)$, a logistic regression fitted on the train slice to predict M1’s errors from the top probability, the margin to the second label and seed disagreement.

M2: uncertainty router. M2 defers the fraction δ of items with the highest $u(x)$. Its threshold is the $(1 - \delta)$ quantile of u on the train slice, so the deferral rate carries over to unseen data. Items not deferred keep M1’s answer, so LLM calls per item equal δ . Each $\delta \in \{0, 0.1, 0.2, 0.3, 0.5, 1\}$ gives one point on an accuracy-vs-cost curve, and the operating point is fixed on the train slice.

M3: candidate-constrained LLM adjudicator. An open-weight instruction-tuned LLM, served locally with vLLM, reads the sub-question, the full question and the answer, and may only choose

among the labels in $C(x)$. Its prompt gives definitions of the candidate labels, a confusion guide with contrast rules for our three largest error pairs (§6), and up to ten boundary examples, the training items nearest the answer in embedding space whose gold labels are in $C(x)$. Without writing out any reasoning, the model scores the tokens of each candidate label, which gives a distribution $q(c | x)$ over $C(x)$ from a single forward pass with almost no output tokens. M3 then combines the two models as $\lambda \log \bar{p} + (1 - \lambda) \log q - \tau \log \pi$ over $C(x)$, with λ and τ fitted by nested cross-validation and prompts designed on the train slice only. These components carried the top systems (§3), and on the winner’s final stage definitions with a confusion guide raised Subtask 2 from 0.491 to 0.593 and ten boundary examples took it to 0.617 (Shi et al., 2026).

M4: distillation and compression. We run M3 once on every training item, using cross-fitted candidates from M1, and mix its distribution q with the gold label as a soft target for retraining M1 (Hinton et al., 2015), with the mixing weight fitted on the train slice. A sharper M1 should need fewer deferrals for the same accuracy (distillation raised the winning team’s 7B model from 0.587 to 0.758 on Subtask 1 but only from 0.495 to 0.509 on Subtask 2, Shi et al., 2026). Finally, M1 and M3 run with 4-bit weights, and each operating point reports the measured calls, tokens and GPU-seconds per item.

Limitations

Annotator disagreement and label reliability.

The labels themselves are not fully reliable. The dataset’s annotators agree only moderately, with a Fleiss’ κ of 0.644 for the clarity level and 0.48 for the nine evasion types (Thomas et al., 2024), and the hardest pairs reach only 0.43 (*General* vs. *Implicit*) and 0.56 (*General* vs. *Deflection*) (Thomas et al., 2026). On our development set only 125 of the 308 items are unanimous, while 150 split two to one and 33 receive three different labels (§2). This affects both training and evaluation. Every training item carries a single label from one annotator, so the models learn one reading of answers that other annotators would label differently, and they cannot learn how disputed an item is. Multi-reference scoring credits a prediction that matches *any* annotator, so a “correct” prediction may be one that most annotators rejected. Scores also depend on how many annotators an item has, and rescored with two annotators, as on the test set, our 8B system drops from 0.543 to 0.524. The task has a low ceiling as well, since a human annotator scored against the other two reaches only 0.643–0.766 (mean 0.684), so differences near that level fall within human disagreement. A predicted label such as *Dodging* is therefore one plausible reading of an answer, not a fact about it or a verdict on the speaker, and it should be reported with its uncertainty.

Development-set evaluation. All results are on the 308-item development set, because the test labels were never released and the evaluation server has closed. The development set has three annotators per item and the test set two. We rescore with every two-annotator subset to approximate the test setting, but we cannot measure the difference in answer length between the splits. With only 308 items, differences between systems carry wide intervals. The gain of the 8B system has a 95% bootstrap interval of [0.054, 0.224], and the all-data variant could not be separated from its control.

Seeds and selection. Seed variance is large (the standard deviation of single-model Subtask 2 scores is 0.043 for the 8B classifier), which is why our claims rest on ten paired seeds. The twelve-epoch schedule, however, was screened on three seeds only, so it is adopted but not yet established. Epochs are chosen on a 345-row slice of the training data, which is itself noisy. The one decision-rule parameter is fitted by nested cross-validation

on the development set, which keeps the score held out but means that the system uses development labels to fit one scalar.

Scope of the comparison. Changing the backbone also changes the training recipe that comes with it (learning rate, epochs and parameter-efficient fine-tuning), so “knowledge” here is shorthand for everything a larger pretrained model brings. We cannot rule out that the pretraining data of the 8B model includes the public interview transcripts the dataset is drawn from. The published comparison points come from other teams’ papers and protocols, which we did not control.

Coverage. None of our models ever predicts the rarest class (*Partial/half-answer*, 14 development reference sets), and the study covers one dataset in one language (English, U.S. presidential interviews).

Ethical Considerations

Our classifiers label the form of an answer, not its truthfulness or the speaker’s intent, so their output should not be used as a judgement of any politician. All data are the public interview transcripts released for the shared task.

Acknowledgments

We thank the organisers of SemEval-2026 Task 6 (CLARITY) for creating the task and the dataset and for running the evaluation.

References

- Pengcheng He, Jianfeng Gao, and Weizhu Chen. 2023. [DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing](#). In *The Eleventh International Conference on Learning Representations (ICLR)*.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. 2015. [Distilling the knowledge in a neural network](#). *Preprint*, arXiv:1503.02531.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. [LoRA: Low-rank adaptation of large language models](#). In *The Tenth International Conference on Learning Representations (ICLR)*.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. 2017. [Focal loss for dense object detection](#). In *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pages 2980–2988.

- Aditya Krishna Menon, Sadeep Jayasumana, Ankit Singh Rawat, Himanshu Jain, Andreas Veit, and Sanjiv Kumar. 2021. [Long-tail learning via logit adjustment](#). In *The Ninth International Conference on Learning Representations (ICLR)*.
- Siva Nuthakki, Sanjay Pulagam, and Sai Woona. 2026. [Team evaluators at SemEval-2026 task 6: Instruction-tuned LLMs for clarity and evasion classification in political interviews](#). In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 3442–3445, San Diego, California, USA. Association for Computational Linguistics.
- Wisarat Peerachaidecho and Attapol Rutherford. 2026. [ChulaNLP at SemEval-2026 task 6: A hybrid BERT-LLM framework for political response clarity and evasion detection](#). In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 1874–1881, San Diego, California, USA. Association for Computational Linguistics.
- Jiawei Ren, Cunjun Yu, Shunan Sheng, Xiao Ma, Haiyu Zhao, Shuai Yi, and Hongsheng Li. 2020. [Balanced meta-softmax for long-tailed visual recognition](#). In *Advances in Neural Information Processing Systems 33 (NeurIPS)*.
- Lingling Shi, Haoyu Jin, Shiquan Wang, Fang Yu, Shuangyong Song, and Xuelong Li. 2026. [TeleAI at SemEval-2026 task 6: A confidence-aware multi-stage reasoning framework with chain-of-thought](#). In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 1591–1609, San Diego, California, USA. Association for Computational Linguistics.
- Taleef Tamsal and Jonathan Rusert. 2026. [PFW at SemEval-2026 task 6: Multi-seed DeBERTa ensembles for political response clarity and evasion classification](#). In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 1518–1525, San Diego, California, USA. Association for Computational Linguistics.
- Amruth Tetakali and Lavanya Tetakali. 2026. [SilkPeak at SemEval-2026 task 6: When politicians dodge — unmasking evasion in political interviews through joint multi-task transformer learning](#). In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 2020–2025, San Diego, California, USA. Association for Computational Linguistics.
- Konstantinos Thomas, Giorgos Filandrianos, Maria Lymperaiou, Chrysoula Zerva, and Giorgos Stamou. 2024. [“I never said that”: A dataset, taxonomy and baselines on response clarity classification](#). In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 5204–5233, Miami, Florida, USA. Association for Computational Linguistics.
- Konstantinos Thomas, Giorgos Filandrianos, Maria Lymperaiou, Chrysoula Zerva, and Giorgos Stamou. 2026. [SemEval-2026 task 6: CLARITY – unmasking political question evasions](#). In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 3704–3715, San Diego, California, USA. Association for Computational Linguistics.
- Tai Tran Tan and An Dinh. 2026. [ttda704 at SemEval-2026 task 6: Structured chain-of-thought prompting for political evasion detection](#). In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 2817–2829, San Diego, California, USA. Association for Computational Linguistics.
- Nawar Turk, Lucas Miquet-Westphal, and Leila Kosseim. 2026. [CLaC at SemEval-2026 task 6: Response clarity detection in political discourse](#). In *Proceedings of the 20th International Workshop on Semantic Evaluation (2026)*, pages 1890–1902, San Diego, California, USA. Association for Computational Linguistics.
- Alexandra N. Uma, Tommaso Fornaciari, Dirk Hovy, Silviu Paun, Barbara Plank, and Massimo Poesio. 2021. [Learning from disagreement: A survey](#). *Journal of Artificial Intelligence Research*, 72:1385–1470.
- An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, et al. 2025. [Qwen3 technical report](#). *Preprint*, arXiv:2505.09388.

A Implementation Details

Encoder. Table 6 lists the DeBERTa settings. Baseline and final models differ only in the input and the number of epochs. One detail mattered: transformers 5.x loads a checkpoint in its stored dtype, and DeBERTa-v3’s published weights are fp16, in which its attention overflows. The model then trained smoothly to the label prior and predicted Explicit for every item, so we load the weights in fp32 explicitly.

Setting	Value
Backbone	microsoft/deberta-v3-large, all weights trained
Head	[CLS] $\rightarrow$ dense + tanh $\rightarrow$ dropout 0.1 $\rightarrow$ linear (9)
Input, baseline	sub-question (≤ 96 tokens) [SEP] answer, 512 tokens
Input, final	“Sub-question: ... Full question: ...” (≤ 256 tokens) [SEP] answer, 1,024 tokens
Truncation	answer cut from the right
Loss	cross-entropy, no smoothing, no class weights
Optimiser	AdamW, $\beta = (0.9, 0.98)$, $\epsilon = 10^{-6}$, weight decay 0.01 (none on biases and LayerNorm)
Learning rate	$10^{-5} \times 0.95^k$ for a layer k levels below the head, head and pooler 10^{-4}
Schedule	cosine, warm-up 10% of steps
Batch	16, no accumulation, gradient-norm clipping 1.0
Epochs	8 (baseline) or 16
Precision	fp32 weights, bf16 autocast, gradient checkpointing
Epoch selection	best macro-F1 on the 345-row train slice (last epoch for all-of-train runs)
Seeds	0–9 (all of train: 0–4)
Logit adjustment	$\tau \in [-1, 2]$, 61-point grid, prior = train label frequencies, 5-fold nested CV on dev
Hardware	one shared RTX PRO 6000 (96 GB), peak memory about 9 GB at 512 tokens, 15–17 GB at 1,024

Table 6: DeBERTa-v3-large training and decision settings.

LLM classifier. Table 7 lists the Qwen3-8B-Base settings, and everything not listed matches the final encoder configuration. Each run used one NVIDIA RTX PRO 6000 (96 GB) on a rented virtual machine. A 3-epoch seed took 15–21 min and a 12-epoch seed about 59 min, and the ten-seed round cost INR 1,518.79 of GPU time.

Code and records. Code, experiment definitions, the experiment log with every pre-registered hypothesis, and the raw analysis outputs behind every number are at <https://github.com/IamSid44/>

Setting	Value
Backbone	Qwen3-8B-Base, bf16
Adapter	LoRA, $r = 16$, $\alpha = 32$, dropout 0.05
LoRA targets	q, k, v, o, gate, up, down_proj
Head	new 9-way linear, last non-pad token, trained in full
Precision	fp32 trainable weights, bf16 autocast
Optimizer	AdamW, lr 10^{-4} , weight decay 0
Gradient clipping	norm 1.0
Schedule	cosine, linear warmup 5% of steps
Epochs	3 (12 in the follow-up)
Batch	4×4 accumulation = 16
Steps per epoch	194 (3,103 training rows)
Loss	cross-entropy, unweighted
Input	question part ≤ 256 tokens, total $\leq 1,024$ incl. closing question
Batching	length-sorted within chunks of 50 batches, order fixed by (seed, epoch)
Epoch selection	best macro-F1 on the 345-row slice (all-of-train runs: last epoch)
Padding check	abort if relative logit difference > 0.2
Seeds	0–9 (12 epochs: 0–2)

Table 7: Hyperparameters of the LLM classifier.

CLARITY-Unmasking-Political-Question-Evasions. Per-seed probabilities, metrics and LoRA adapters are at <https://huggingface.co/siddharthg44/clarity-semeval26>, and training curves at <https://wandb.ai/iamsid44-iiit-hyderabad/clarity-semeval26>. Software: Python 3.12, PyTorch 2.11, Transformers 5.17, PEFT 0.20.