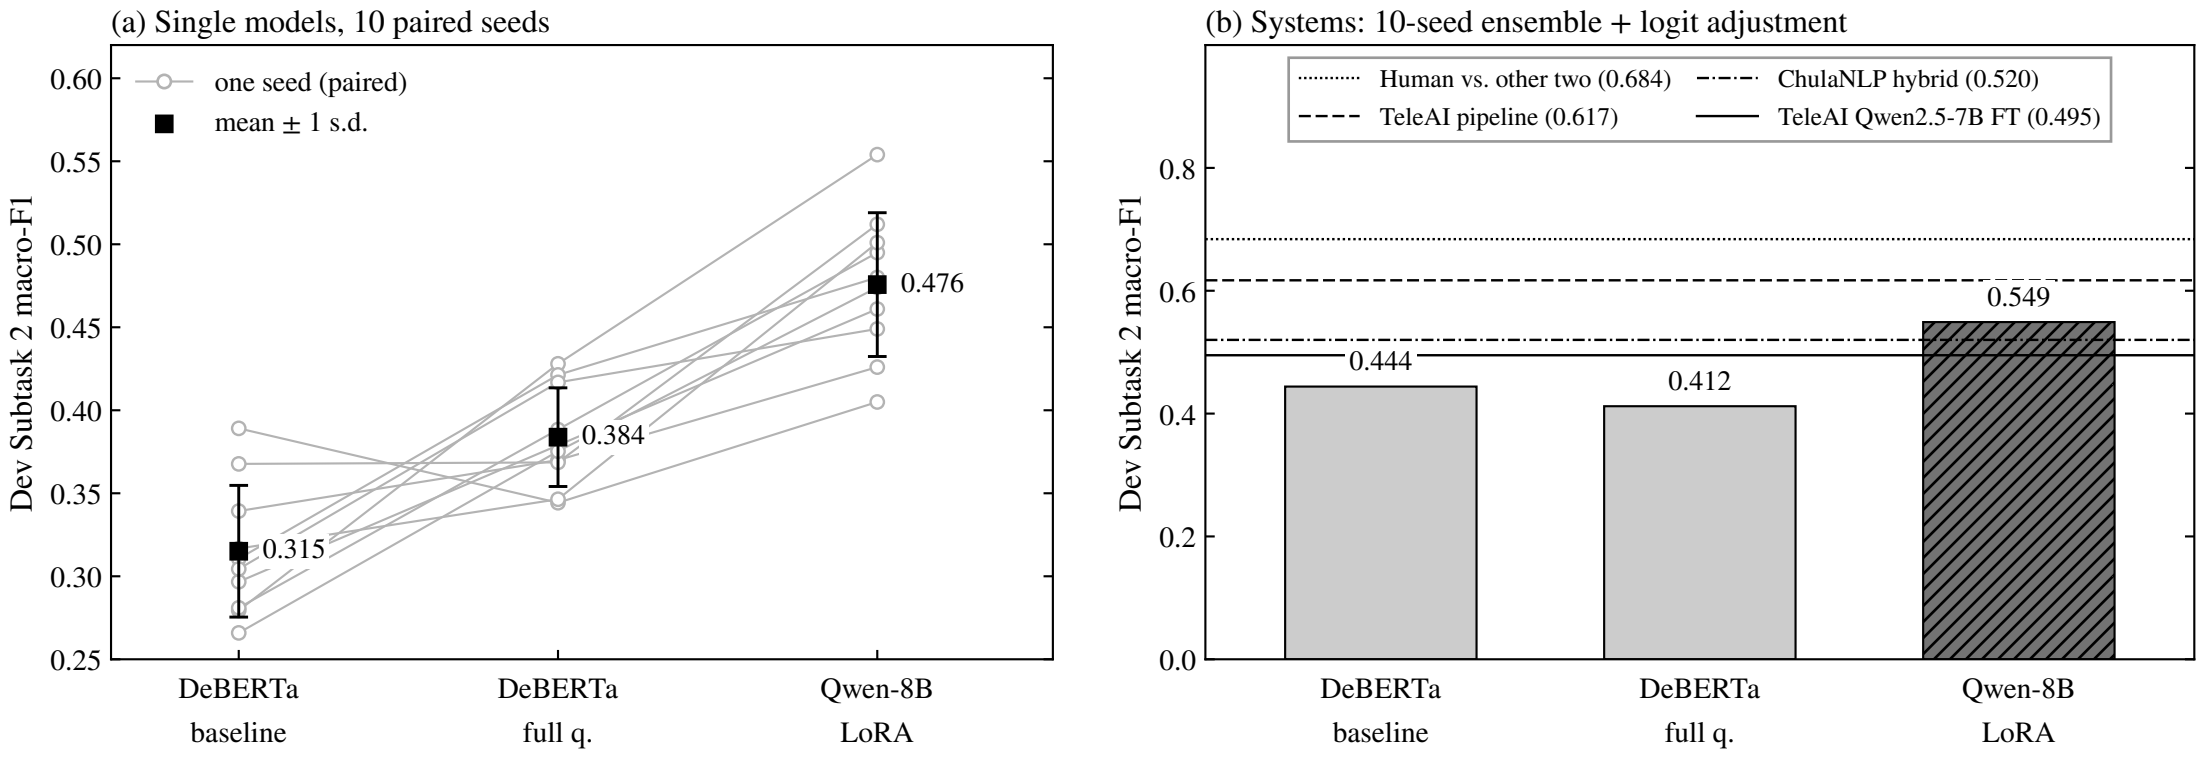
(a) Single models, 10 paired seeds
one seed (paired)
mean ± 1 s.d.
Dev Subtask 2 macro-F1
0.60
0.55
0.50
0.45
0.40
0.35
0.30
0.25
0.315
0.384
0.476
DeBERTa baseline
DeBERTa full q.
Qwen-8B LoRA
(b) Systems: 10-seed ensemble + logit adjustment
Human vs. other two (0.684)
ChulaNLP hybrid (0.520)
TeleAI pipeline (0.617)
TeleAI Qwen2.5-7B FT (0.495)
Dev Subtask 2 macro-F1
0.8
0.6
0.4
0.2
0.0
0.444
0.412
0.549
DeBERTa baseline
DeBERTa full q.
Qwen-8B LoRA